

Understanding k8s workloads



THAO LUONG
03/2022

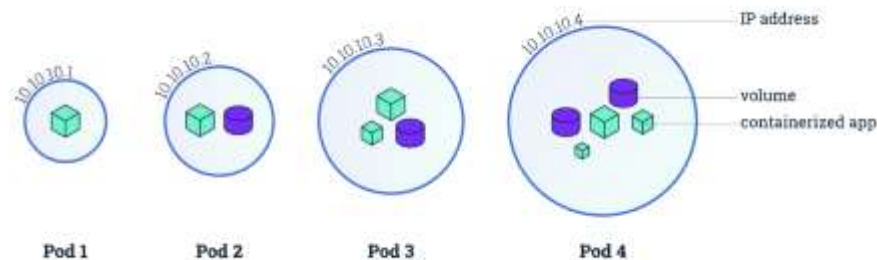


Content

- ❑ Understanding Pods, Nodes
- ❑ Controllers and Pods
- ❑ Deployment , ReplicaSet
- ❑ Multi-container Pods
- ❑ Managing Pod Health with Probes



What is a Pod?

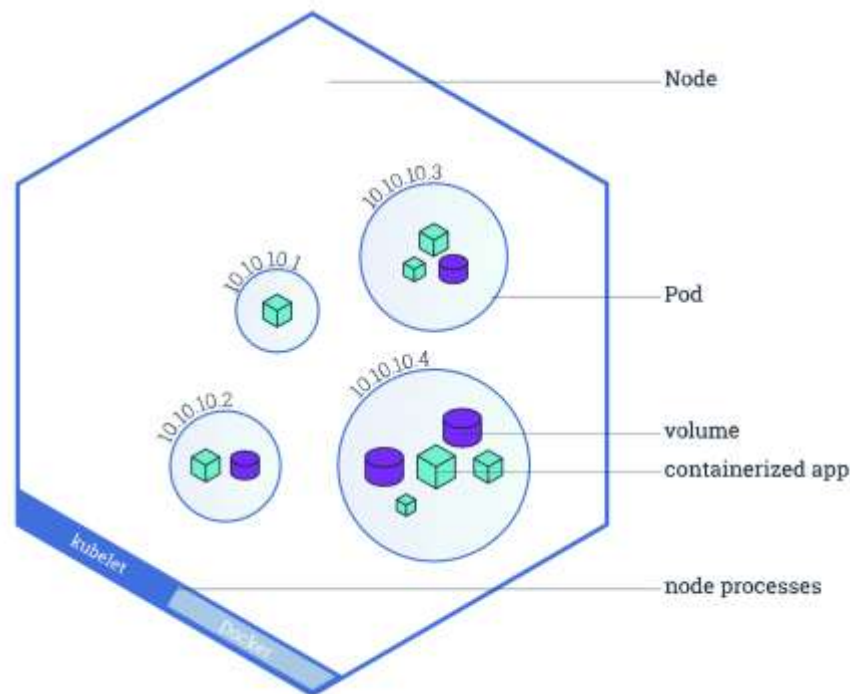


- A group of whales called “pod” = a group of one or more containers
- Shared resources like networking, storages and specification (image version, container ports...)
- “Logical host”
- Atomic unit on the k8s platform



Pods & Nodes

- Node can be a VM, bare-metal or even Docker container (kubernetes in Docker)
- Node runs at a least: kubelet + container runtime
- Node failure -> **identical Pods** are scheduled on other nodes.





How pods manage containers



Single Container Pods



Multi-Container Pods



Single Pod manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world-pod
spec:
  containers:
  - name: hello-world
    image: gcr.io/google-samples/hello-app:1.0
    ports:
    - containerPort: 80
```



Controllers & Pods



Controllers keep your apps in the desired state

Responsible for starting and stopping Pods

Application scaling

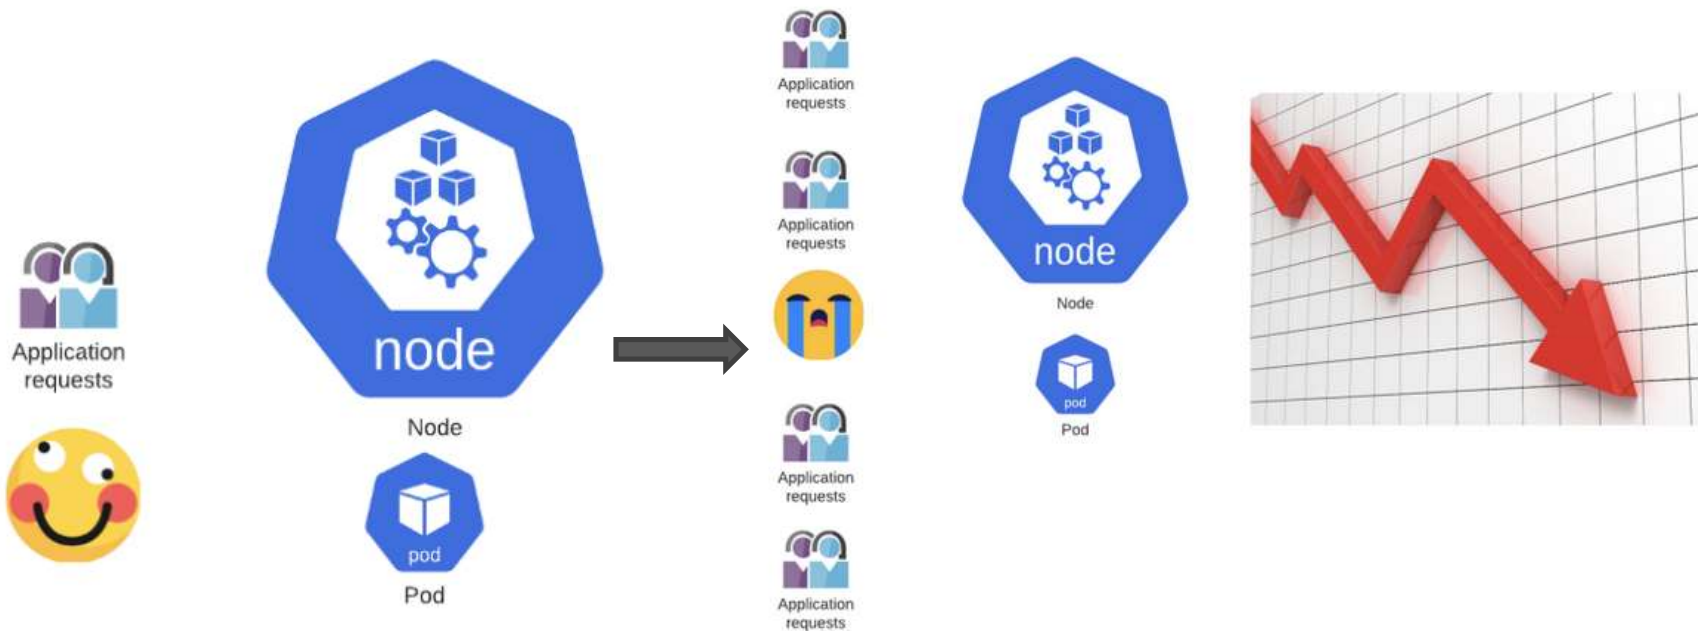
Application recovery

You don't want to run bare/naked Pods

They won't be recreated in the event of a failure



Understanding ReplicaSets

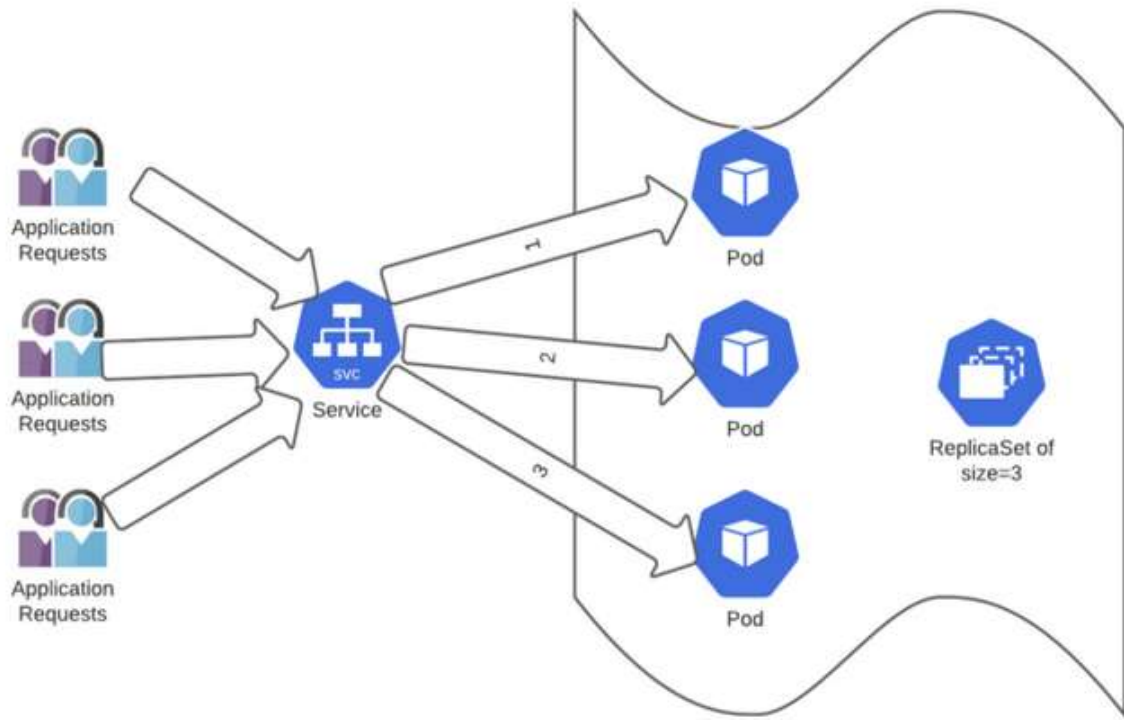


Single POD application

Single pod with increased application requests seem to bring a lot of sadness.



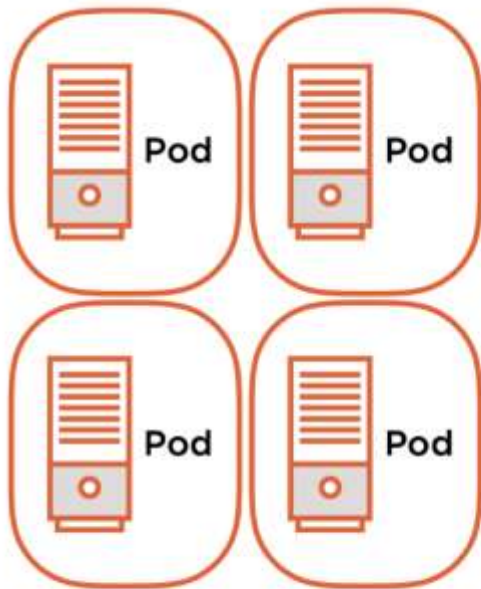
Understanding ReplicaSets



User requests getting routed to application PODs through the K8s service object.



Understanding ReplicaSets



Deploys a defined number of Pods

Consists of a Selector, Number of Replicas (Pods) and a Pod Template

Generally speaking you don't create ReplicaSets directly

You create Deployments



Managing app state with deployment



Creating



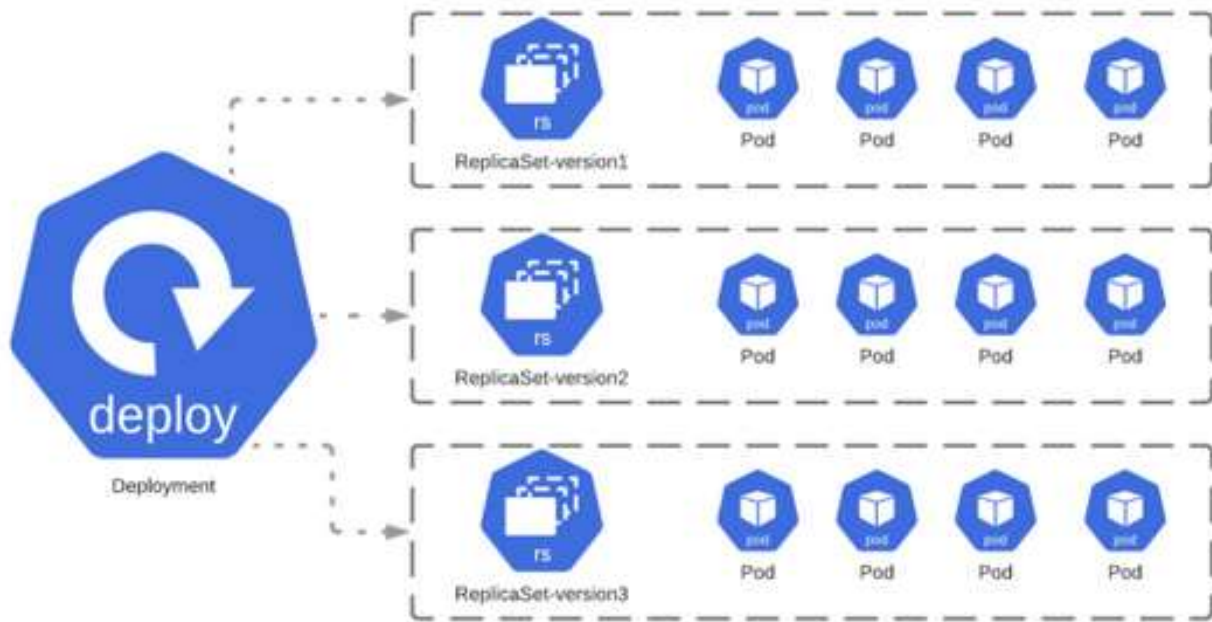
Updating



Scaling



Managing app state with deployment



The relation between Deployment, Replicaset & POD



Pod/ReplicaSet/Deployment

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: hello-world
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello-world-pod
  template:
    metadata:
      labels:
        app: hello-world-pod
    spec:
      containers:
        - name: hello-world
          image: gcr.io/google-samples/hello-app:1.0
```

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world-pod
spec:
  containers:
    - name: hello-world
      image: gcr.io/google-samples/hello-app:1.0
      ports:
        - containerPort: 8080
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-world-1
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello-world-1
  template:
    metadata:
      labels:
        app: hello-world-1
    spec:
      containers:
        - name: hello-world
          image: gcr.io/google-samples/hello-app:1.0
          ports:
            - containerPort: 8080
```



Multi-container Pods



Tightly coupled applications

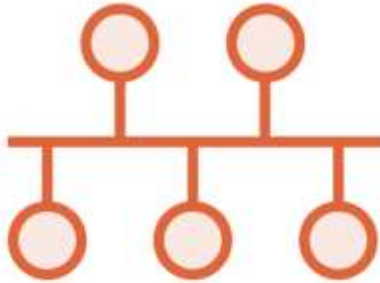
Scheduling processes together

Requirement on some shared resource

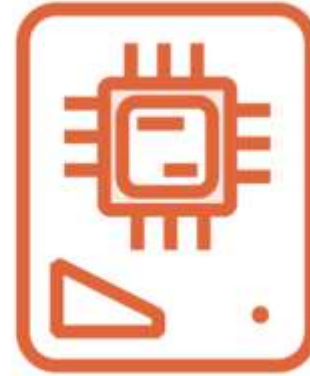
Usually something generating data while the other process consumes



Share resources inside a pod



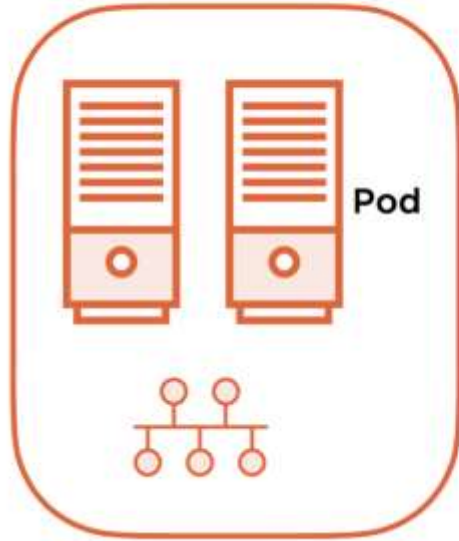
Networking



Storage



Shared networking



Shared loopback interface, used for communication over localhost

Be mindful of application port conflicts



Shared storage



Each container image has it's own file system

Volumes are defined at the Pod level

Shared amongst the containers in a Pod

Mounted into the containers' file system

Common way for containers to exchange data

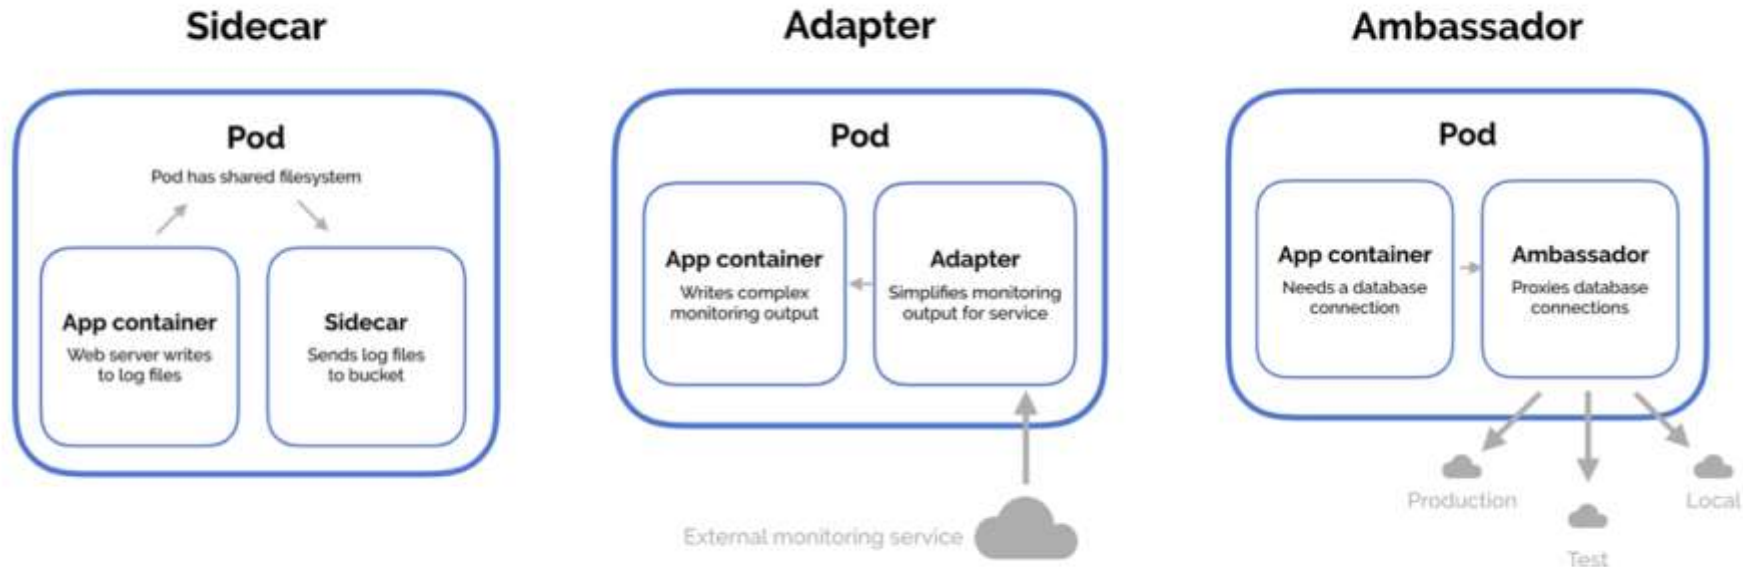


Shared storage

```
apiVersion: v1
kind: Pod
metadata:
  name: mc1
spec:
  volumes:
  - name: html
    emptyDir: {}
  containers:
  - name: 1st
    image: nginx
    volumeMounts:
    - name: html
      mountPath: /usr/share/nginx/html
  - name: 2nd
    image: debian
    volumeMounts:
    - name: html
      mountPath: /html
    command: ["/bin/sh", "-c"]
    args:
    - while true; do
      date >> /html/index.html;
      sleep 1;
    done
```



Common multi-container pods design





Terminating a pod

Grace Period Timer
(30 sec default)

Pods changes to
Terminating

SIGTERM

Service Endpoints
and Controllers
updated

IF > Grace Period
SIGKILL

API and etcd are
updated

```
kubectl delete pod --grace-period=<seconds>
```

Force Deletion - Immediately deletes records in API and etcd



Container Restart Policy



A container in a Pod can restart independent of the Pod

Applies to containers inside a Pod and defined inside the Pod's Spec

The Pod is the environment the container runs in

Not rescheduled to another Node, but restarted by the Kubelet on that Node

Restarts with an exponential backoff, 10s, 20s, 40s capped at 5m and reset to 0 after 10m of successful runtime



Container Restart Policy



Always (default) - will restart all containers inside a Pod

OnFailure - Non-graceful termination

Never



Defining pod health



A Pod is considered ready when all containers are ready

But we'd like to be able to understand a little more about our applications

We can add additional intelligence to our Pod's state and health

Container Probes

`livenessProbe`

`readinessProbe`



livenessProbes



Runs a diagnostic check on a container

Per container setting

On failure, the Kubelet restarts the container

Container Restart Policy

Give Kubernetes a better understanding of our application



readinessProbes



Runs a diagnostic check on the container

Per container setting

On startup, your application won't receive traffic until ready

On failure, removes Pod from load balancing or replication controller

Applications have long startup times

Prevents users from seeing errors



Type of checks

Exec	tcpSocket	httpGet
Process exit code	Successfully Open a Port	Return Code 200 => and < 400
Success	Failure	Unknown



Configuring probes



`initialDelaySeconds` - number of seconds after the container has started before running container probes

`periodSeconds` - probe interval, default 10 seconds

`timeoutSeconds` Probe timeout 1 seconds

`failureThreshold` - number of missed checks before reporting failure, default 3

`successThreshold` - number of probes to be considered successful and live, default 1

