

Application Life Cycle management



THAO LUONG
09/2020

<https://medium.com/@luongvinhthao>



Content

- ❑ Rolling update and Rollback deployment
- ❑ Managing and Controlling the Kubernetes Schedule



Deploying Application with Deployment



Creating



Updating



Scaling



Update deployment object

```
kubectl set image deployment hello-world hello-world=hello-app:2.0
```

```
kubectl set image deployment hello-world hello-world=hello-app:2.0 --record
```

```
kubectl edit deployment hello-world
```

```
kubectl apply -f hello-world-deployment.yaml --record
```



Checking deployment status



```
kubectl rollout status deployment [name]
```

```
kubectl describe deployment [name]
```

Deployment Status

Complete - all update work is finished

Progressing - update in flight

Failed - update could not complete



Using Deployments to Change State



Control rollouts of a new version of your application

Update Strategy

Pause to make corrections

Rollback to an earlier version



Controlling Rollouts With Update Strategy



Controls Pods rollout

RollingUpdate (Default)

A new ReplicaSet starts scaling up and the old ReplicaSet starts scaling down

Recreate

Terminates all Pods in the current ReplicaSet set prior to scaling up the new ReplicaSet

Used when applications don't support running different versions concurrently



Controlling Deployment Rollouts



Update Strategy in a Deployment Spec

Readiness Probes in your Pod Template Spec

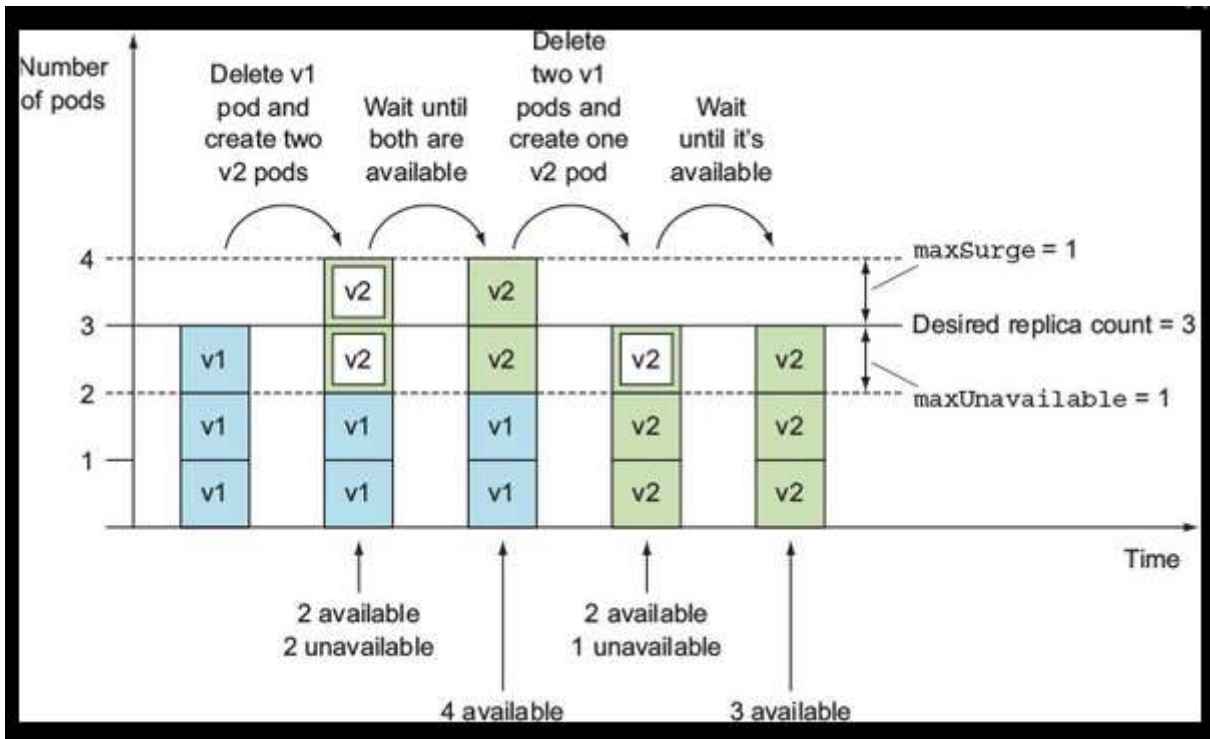


Update Strategy

```
apiVersion: apps/v1
kind: Deployment
...
spec:
  replicas: 20
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 20%
      maxSurge: 5
  ...
  template:
    ...
    spec:
      containers:
        ...
        readinessProbe:
          httpGet:
            path: /index.html
            port: 8080
            initialDelaySeconds: 10
            periodSeconds: 10
```




Rolling update with maxSurge = 1 and maxUnavailable = 1





Pausing and Resuming a Deployment



Changes to the Deployment while paused are not rolled out

Batch changes together, then resume the rollout

The current state of the Deployment is maintained until it's resumed

Starts up a new ReplicaSet with the new changes

```
kubectl rollout pause deployment \
my-deployment
```

```
kubectl rollout resume my-deployment
```



Rolling Back a Deployment



Rollout history

CHANGE-CAUSE Annotation Deployment

Revision History

revisionHistoryLimit defaults to 10

Number of ReplicaSets retained in history

Used for rolling back

Can be set to 0 for immediate cleanup



Rolling Back a Deployment(cont')



```
kubectl rollout history deployment \
hello-world
```

```
kubectl rollout history deployment \
hello-world --revision=1
```

```
kubectl rollout undo deployment
hello-world
```

```
kubectl rollout undo deployment \
hello-world --to-revision=1
```



Deployment Tips



Control your rollouts with an Update Strategy appropriate for your application




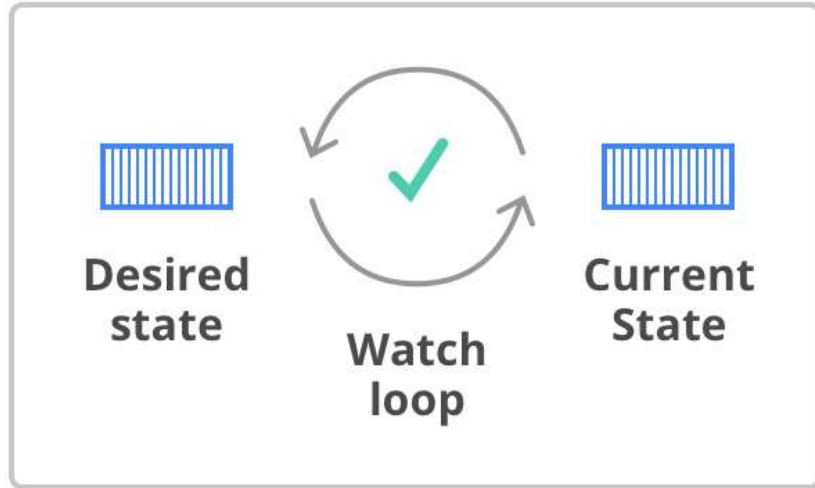
Use Readiness Probes for your application



Use the `--record` option to leave a trail of your work for others




Self-Healing concept







Scheduling in Kubernetes



Selecting a Node to start a Pod on
kube-scheduler



Scheduling in Kubernetes



Resources



Policy



Scheduling Process

Watches the API
Server for
Unscheduled Pods

Node selection

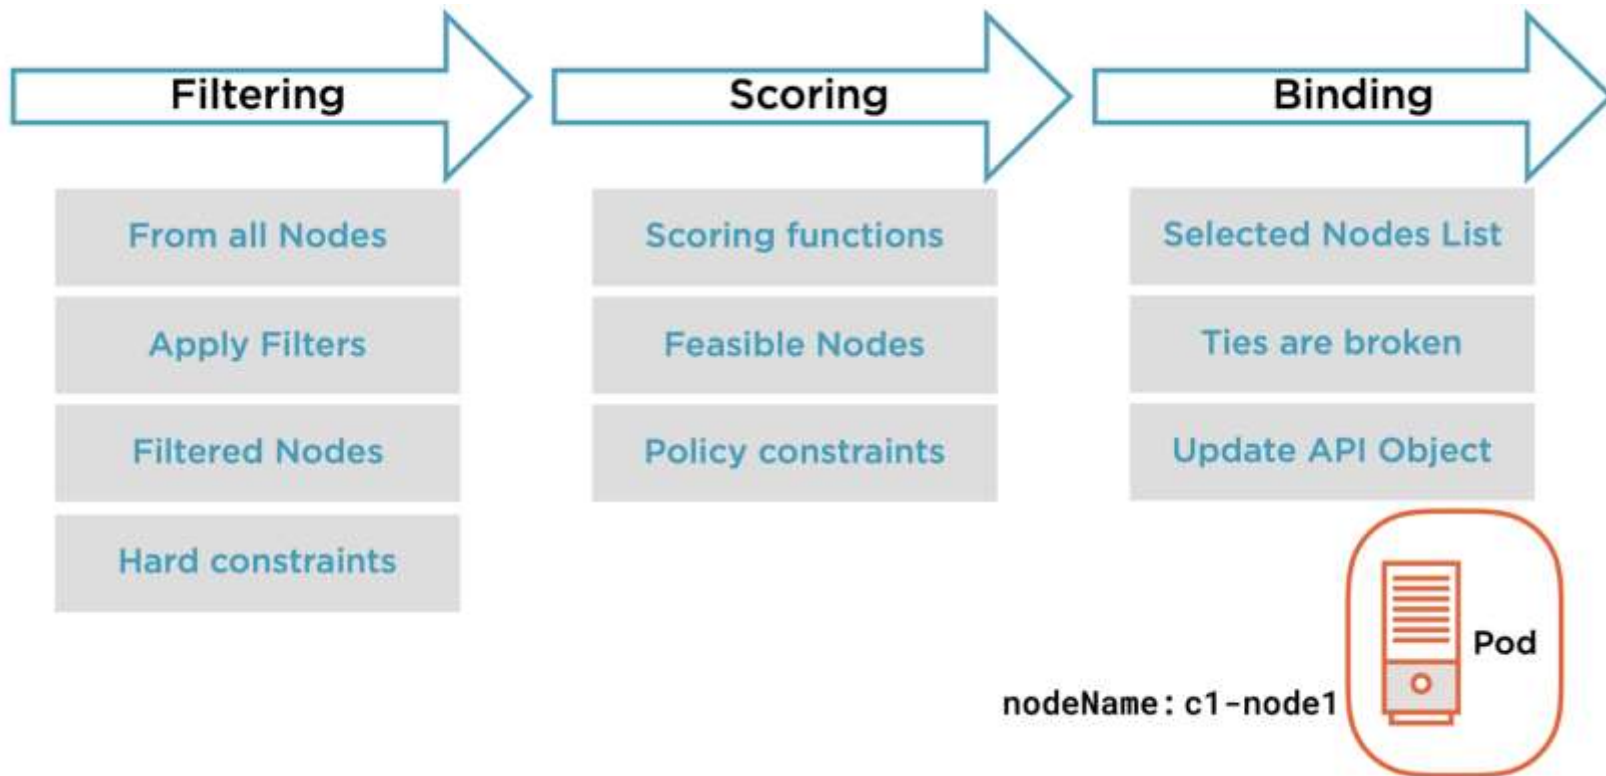
Update nodeName in
the Pod object

Nodes' kubelets
watch API Server
for work

Signal container
runtime to start
container(s)



Node Selection





Resource Requests



Setting requests will cause the scheduler to find a Node to fit the workload/Pod

requests are guarantees

CPU

Memory

Allocatable resources per Node

Pods that need to be scheduled but there not enough resources available will go Pending



Controlling Scheduling

Node Selector

Affinity

Taint and
Tolerations

Node Cordoning

Manual Scheduling



Node Selector



nodeSelector - assign Pods to Nodes using Labels and Selectors

Apply Labels to Nodes

Scheduler will assign Pods to a Node with a matching Label

Simple key/value check based on `matchLabels`

Often used to map Pods to Nodes based on...


- Special hardware requirements

- Workload isolation

Managing Kubernetes API Server and Pods



Scheduling - Node Selector





Assigning Pods to Nodes using Node

```
kubectl label node c1-node3 hardware=local_gpu
```

```
spec:
```

```
  containers:
```

```
  - name: hello-world
```

```
    image: gcr.io/google-samples/hello-app:1.0
```

```
    ports:
```

```
    - containerPort: 8080
```

```
  nodeSelector:
```

```
    hardware: local_gpu
```



Affinity and Anti-Affinity



nodeAffinity - uses Labels on Nodes to make a scheduling decision with matchExpressions

`requiredDuringSchedulingIgnoredDuringExecution`

`preferredDuringSchedulingIgnoredDuringExecution`

podAffinity - schedule Pods onto the same Node, Zone as some other Pod

podAntiAffinity - schedule Pods onto the different Node, Zone as some other Pod



Using Affinity to Control Pod Placement

```
spec:
  containers:
  - name: hello-world-cache
  ...
  affinity:
    podAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
      - labelSelector:
          matchExpressions:
          - key: app
            operator: In
            values:
            - hello-world-web
        topologyKey: "kubernetes.io/hostname"
```



Taints and Tolerations



Taints - ability to control which Pods are scheduled to Nodes

Tolerations - allows a Pod to ignore a Taint and be scheduled as normal on Tainted Nodes


Useful in scenarios where the cluster administrator needs to influence scheduling without depending on the user

key=value:effect

```
kubectl taint nodes c1-node1 \
  key=MyTaint:NoSchedule
```



Scheduling - Taints and Tolerations





Adding a Taint to a Nodes and a Toleration to a Pod

```
kubectl taint nodes c1-node1 key=MyTaint:NoSchedule
```

```
spec:
```

```
  containers:
```

```
  - name: hello-world
```

```
    image: gcr.io/google-samples/hello-app:1.0
```

```
    ports:
```

```
    - containerPort: 8080
```

```
  tolerations:
```

```
  - key: "key"
```

```
    operator: "Equal"
```

```
    value: "MyTaint"
```

```
    effect: "NoSchedule"
```



Node Cordoning



Marks a Node as unschedulable

Prevents new Pods from being scheduled to that Node

Does not affect any existing Pods on the Node

This is useful as a preparatory step before a Node reboot or maintenance

```
kubectl cordon c1-node3
```

If you want to gracefully evict your Pods from a Node...

```
kubectl drain c1-node3 --ignore-daemonsets
```



DeScheduler

- Some nodes are under or over utilized.
- The original scheduling decision does not hold true any more, as taints or labels are added to or removed from nodes, pod/node affinity requirements are not satisfied any more.
- Some nodes failed and their pods moved to other nodes.
- New nodes are added to clusters.

```
apiVersion: "descheduler/v1alpha1"
kind: "DeschedulerPolicy"
strategies:
  "LowNodeUtilization":
    enabled: true
    params:
      nodeResourceUtilizationThresholds:
        thresholds:
          "cpu" : 20
          "memory": 20
          "pods": 20
        targetThresholds:
          "cpu" : 50
          "memory": 50
          "pods": 50
```




DeScheduler

- Some nodes are under or over utilized.
- The original scheduling decision does not hold true any more, as taints or labels are added to or removed from nodes, pod/node affinity requirements are not satisfied any more.
- Some nodes failed and their pods moved to other nodes.
- New nodes are added to clusters.

```
apiVersion: "descheduler/v1alpha1"
kind: "DeschedulerPolicy"
strategies:
  "LowNodeUtilization":
    enabled: true
    params:
      nodeResourceUtilizationThresholds:
        thresholds:
          "cpu" : 20
          "memory": 20
          "pods": 20
        targetThresholds:
          "cpu" : 50
          "memory": 50
          "pods": 50
```

