

Kubernetes networking management



THAO LUONG
09/2020



Content

- ☐ Kubernetes network model
- ☐ Network topology
- ☐ Pod networking
- ☐ Container Network Interface
- ☐ CoreDNS



k8s networking model

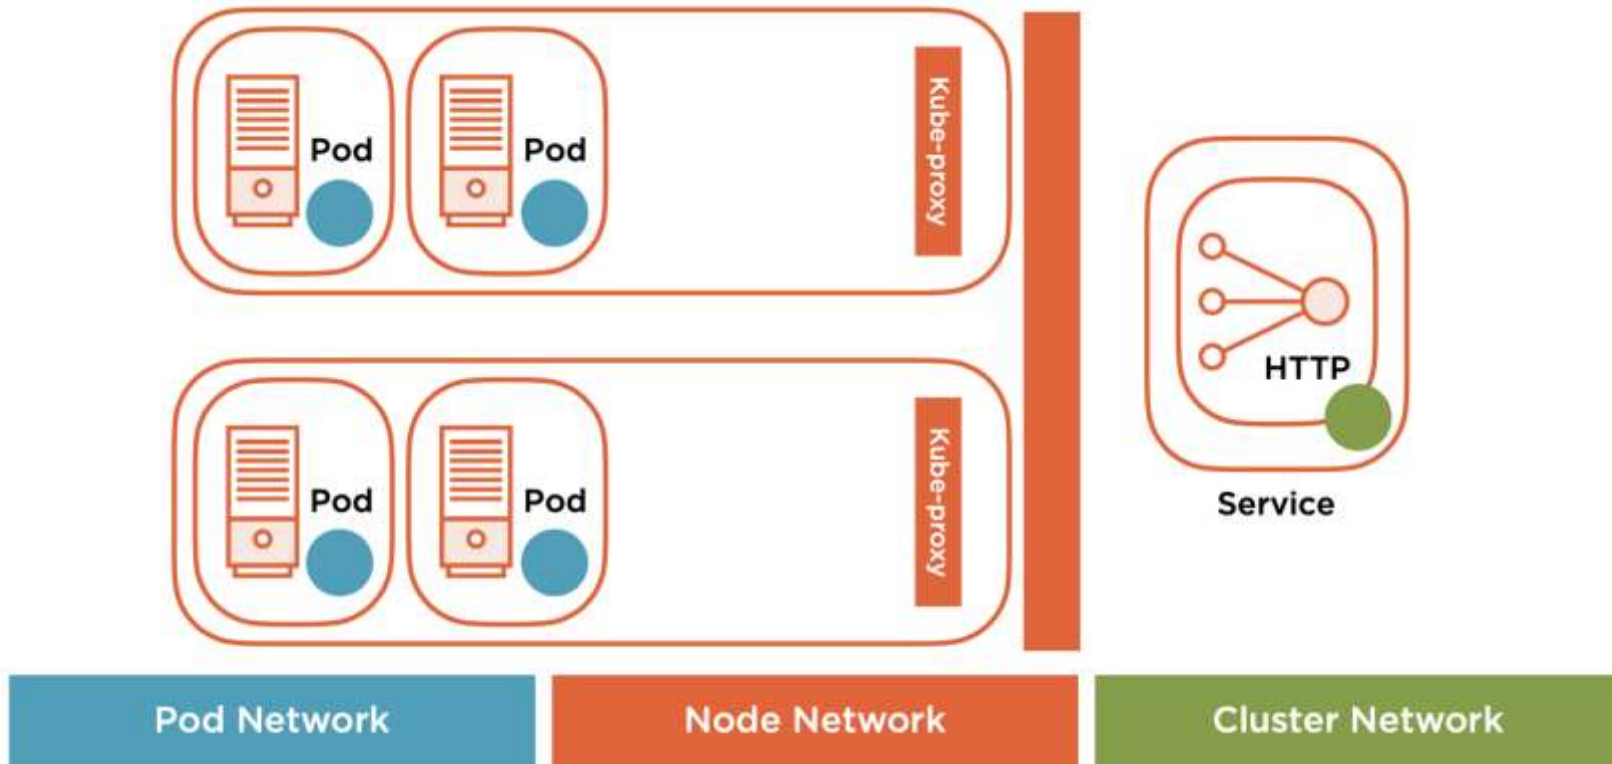
**All Pods can
communicate with
each other on all
Nodes**

**Agents on a Node
can communicate
with all Pods on
that Node**

**No Network Address
Translation (NAT)**

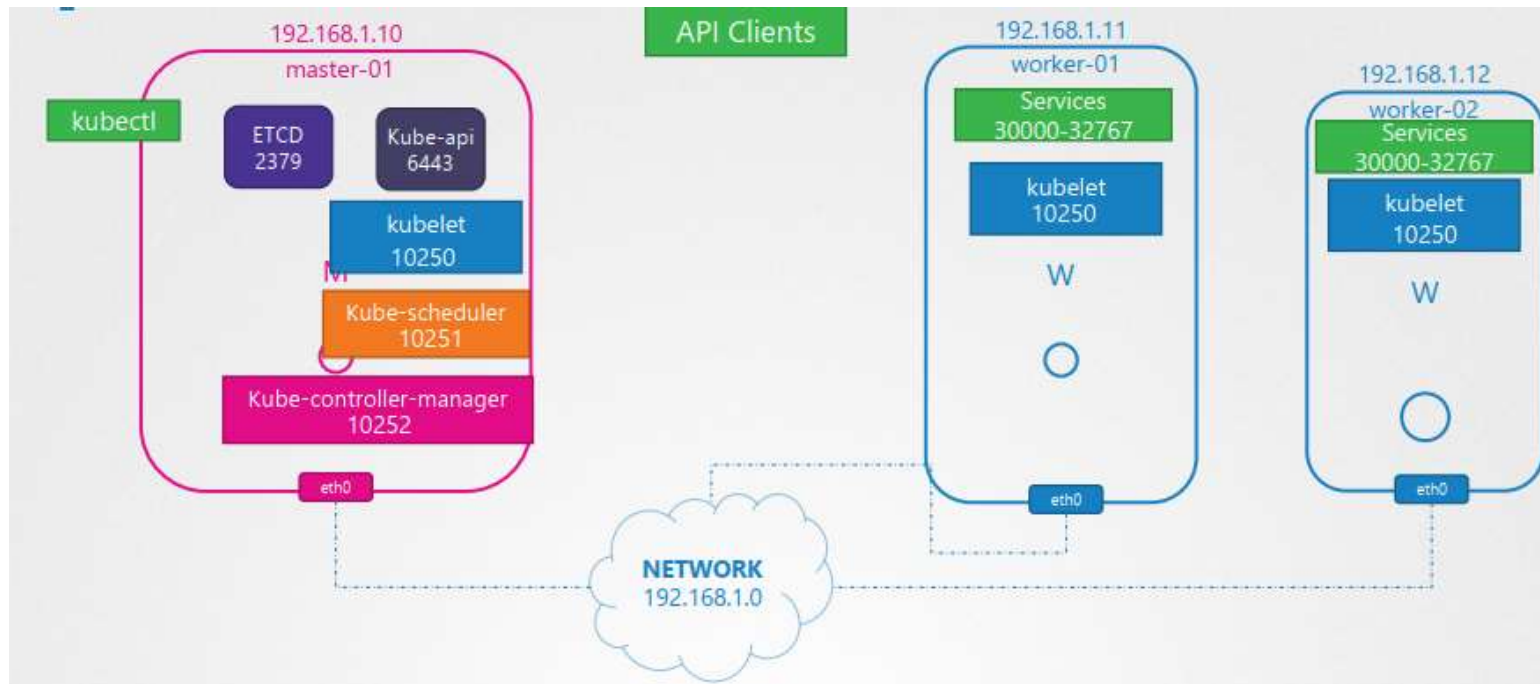


Network Topology





Networking cluster Nodes





Pod Networking



Pod share a network namespace

Containers in a Pod communicate over localhost

Pause/Infrastructure container

Starts the networking namespace

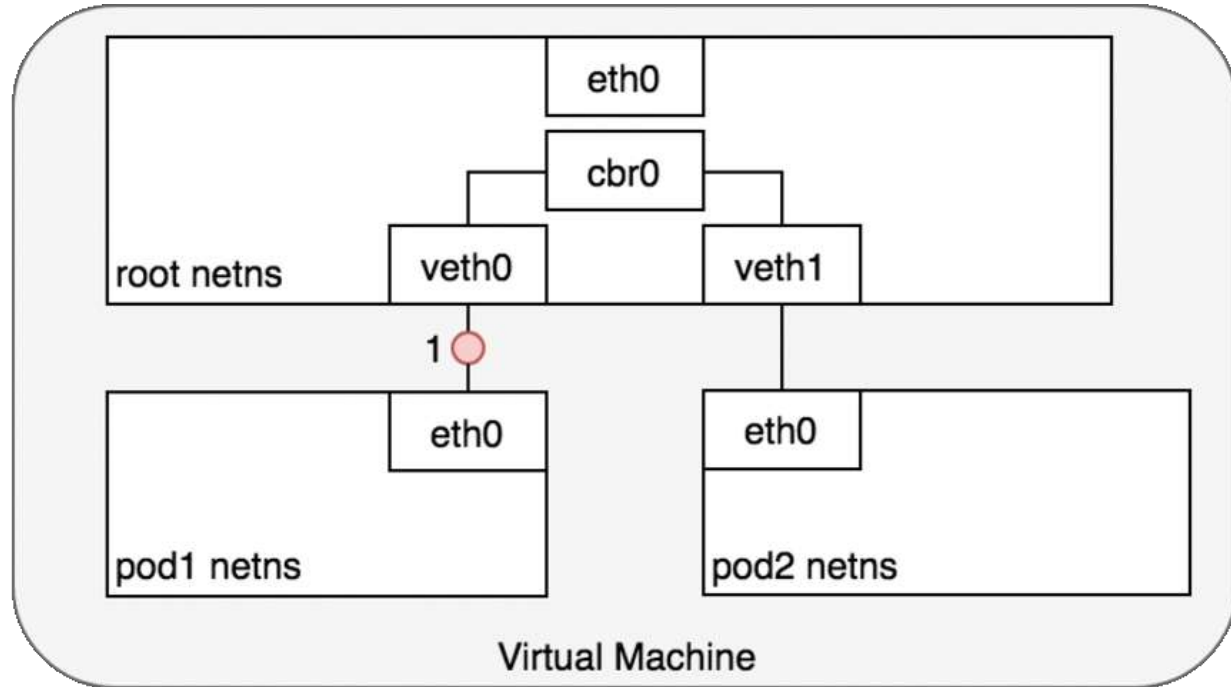
If the application container restarts the network will persist

Lifecycle of the Pod



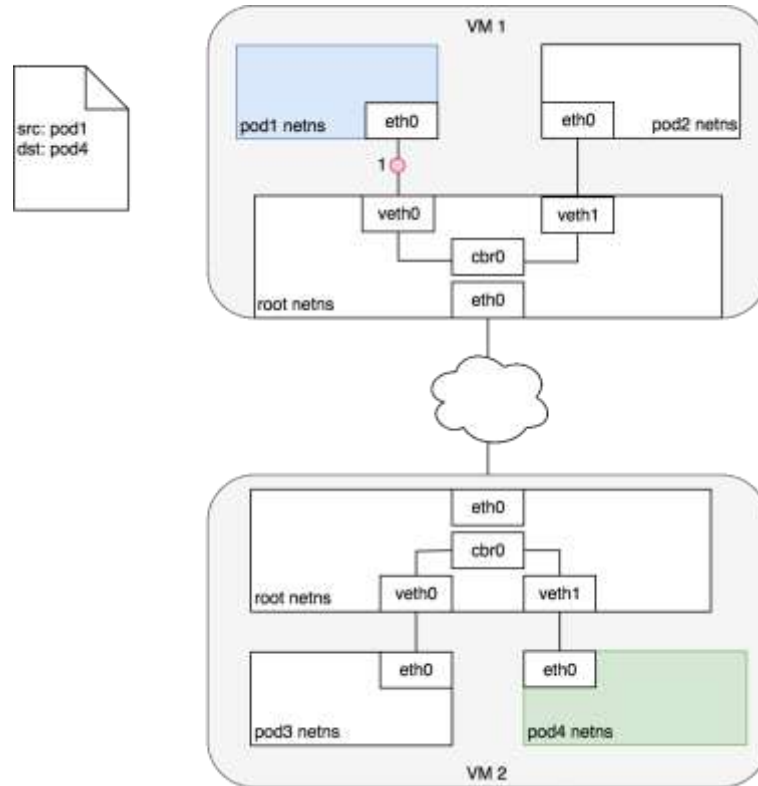
Pod-to-Pod Networking

src: pod1
dst: pod2





Pod-to-Pod Networking Across Nodes





Route Table of Nodes

```
gke-cluster-1-default-pool-487a6374-lj81 / # route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	10.148.0.1	0.0.0.0	UG	1024	0	0	eth0
10.48.2.0	0.0.0.0	255.255.255.0	U	0	0	0	cbr0
10.148.0.1	0.0.0.0	255.255.255.255	UH	1024	0	0	eth0
169.254.123.0	0.0.0.0	255.255.255.0	U	0	0	0	docker0

```
gke-cluster-1-default-pool-487a6374-nz29 / # route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	10.148.0.1	0.0.0.0	UG	1024	0	0	eth0
10.48.1.0	0.0.0.0	255.255.255.0	U	0	0	0	cbr0
10.148.0.1	0.0.0.0	255.255.255.255	UH	1024	0	0	eth0
169.254.123.0	0.0.0.0	255.255.255.0	U	0	0	0	docker0



k8s with Calico

```
gke-cluster-2-calico-default-pool-2bb75e10-0ztm / # route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	10.148.0.1	0.0.0.0	UG	1024	0	0	eth0
10.24.1.2	0.0.0.0	255.255.255.255	UH	0	0	0	caliedfbbfb72f7
10.24.1.3	0.0.0.0	255.255.255.255	UH	0	0	0	califb09e4eac3d
10.24.1.4	0.0.0.0	255.255.255.255	UH	0	0	0	calieecd6fd5bbe
10.148.0.1	0.0.0.0	255.255.255.255	UH	1024	0	0	eth0
169.254.123.0	0.0.0.0	255.255.255.0	U	0	0	0	docker0

```
gke-cluster-2-calico-default-pool-2bb75e10-rk6f / # route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	10.148.0.1	0.0.0.0	UG	1024	0	0	eth0
10.24.2.4	0.0.0.0	255.255.255.255	UH	0	0	0	cali9223a8d59a0
10.24.2.5	0.0.0.0	255.255.255.255	UH	0	0	0	cali3b8b1865e83
10.24.2.6	0.0.0.0	255.255.255.255	UH	0	0	0	caliad6916ee747
10.24.2.7	0.0.0.0	255.255.255.255	UH	0	0	0	calia074902df36
10.148.0.1	0.0.0.0	255.255.255.255	UH	1024	0	0	eth0
169.254.123.0	0.0.0.0	255.255.255.0	U	0	0	0	docker0



CNI



CNI (Container Network Interface) is a project by CNCF that defines a specification which allows communication between containers. Kubernetes supports CNI plugins for the communication between pods



CNI

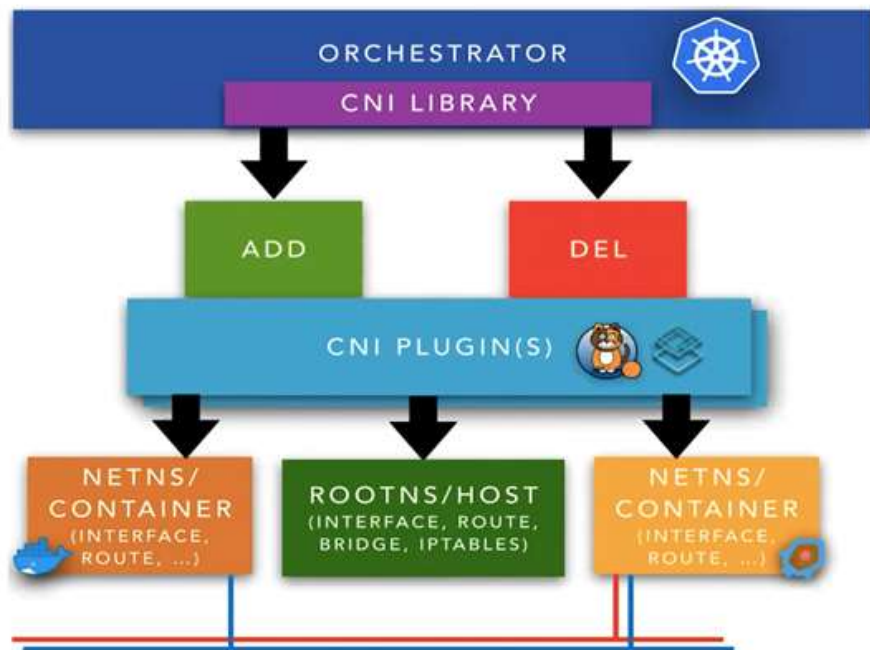


Figure 1: CNI plugin architecture



CNI

- **Calico**
 - one of the most popular plugins
 - default choice of the most of Kubernetes platforms (kubespray, docker enterprise, etc.)
 - supports IP-IP encapsulation if BGP cannot be used
 - supports Network Policies
 - uses iptables for routing but it can be configured to use kube-proxy's IPVS mode
- **Weave**
 - provides VXLAN tunneling solution
 - all of the nodes are connected as mesh which allows it to run on partially connected networks
 - does not scale well because of the mesh structure
 - stores configuration files on pods instead of Kubernetes CRDs or etcd
 - has an encryption library
 - supports Network Policies



Cluster DNS



DNS is available as a Service in a Cluster

Pods are configured to use this DNS

DNS records

- Services - A/AAAA records

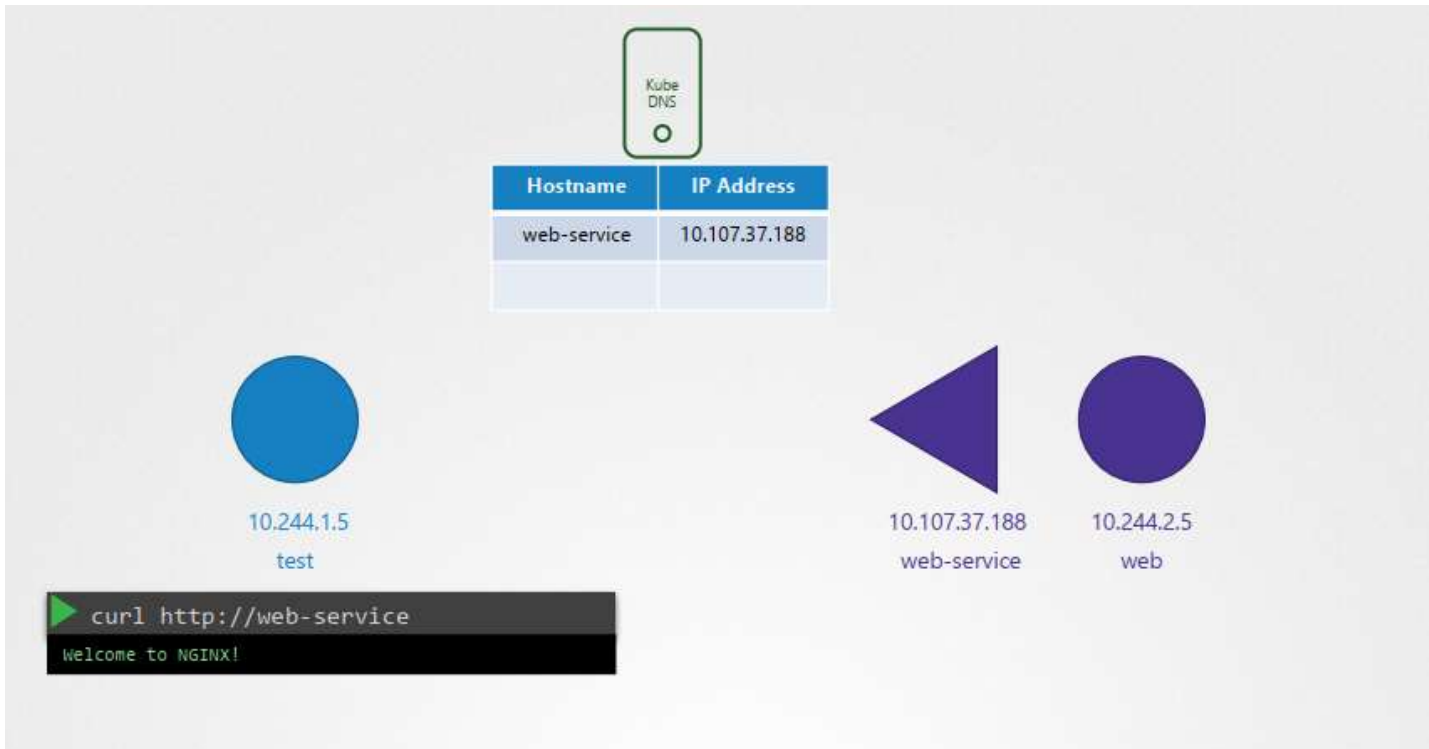
- Namespaces - subdomains

Core to Service discovery

Customize both the DNS Service and Pods configuration



Cluster DNS





Cluster DNS

