

Configuring and managing application access with services



THAO LUONG
06/2020



Content

- ❑ Kube proxy
- ❑ Understanding services
- ❑ Type of service
- ❑ Service network internals
- ❑ Service discovery



Kube-proxy



Kube-proxy reflects service as defined in the Kubernetes API on each node and can do simple TCP, UDP and SCTP stream forwarding or round robin TCP, UDP forwarding across a set of backends.




Kube-proxy

- Iptable mode: Iptable mode from within the kernel, directly forwards the connection to the pod, Fast and efficient but harder to debug
- Ipvs mode when the Client Pod request reaches the kernel space, it is directly distributed to each Pod according to the ipvs rule



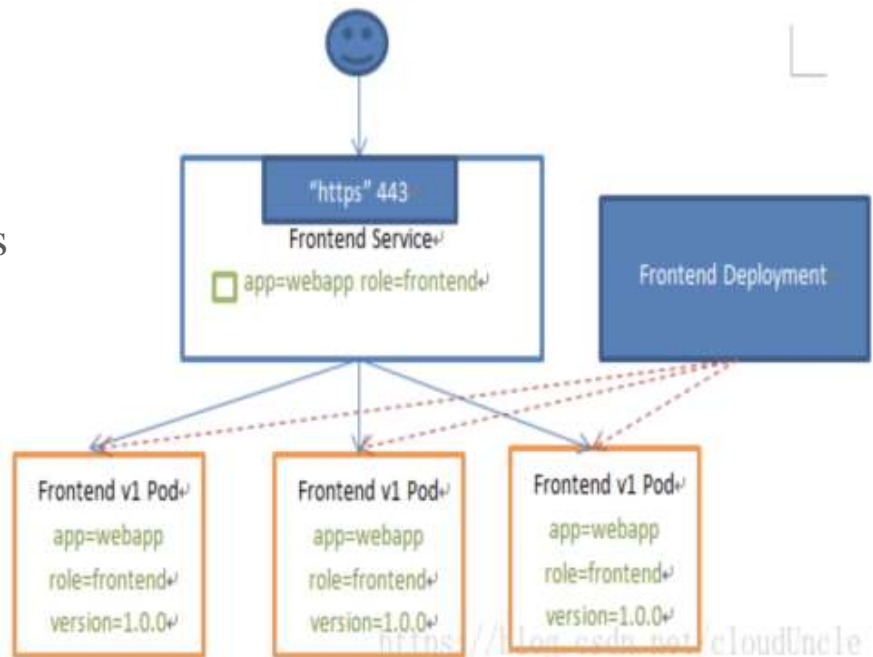
Services





Understanding services

- Service in K8s is an object, similar to a Pod
- A Service can be defined as a logical set of pods.





How services work



Services match Pods using Labels and Selectors

Creates and registers Endpoints in the Service (Pod IP and Port pair)

Implemented in the kube-proxy on the Node in iptables

kube-proxy watches the API Server and the Endpoints

Managing the Kubernetes API Server and Pods



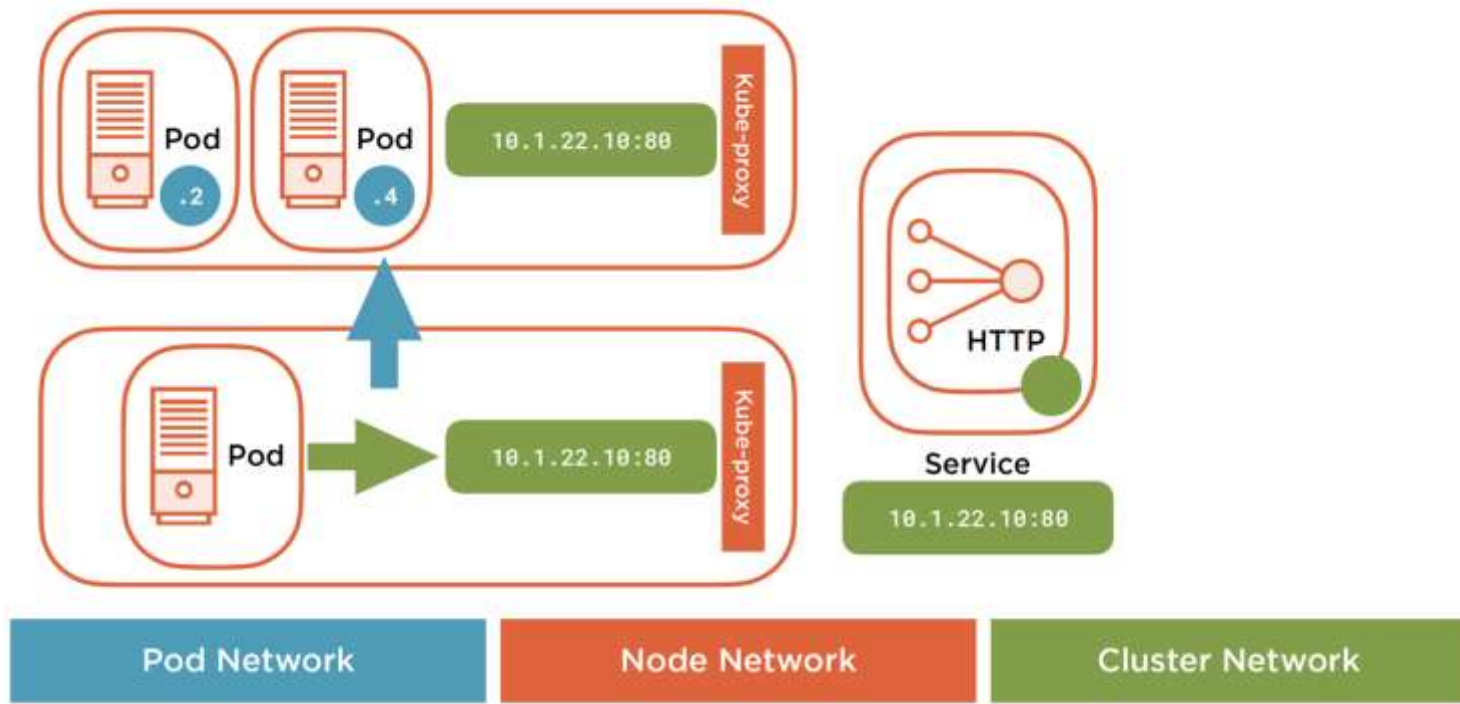
Service types



- **ClusterIp:** The default type, which automatically assigns a virtual IP that can be accessed only inside the Cluster.
- **NodePort:** Bind a port to each machine for Service on the basis of ClusterIP, so you can access the service through <NodeIP>: NodePort
- **LoadBalancer:** Based on the NodePort, create an external load balancer with the cloud provider and forward the request to <NodeIP>: NodePort



ClusterIP

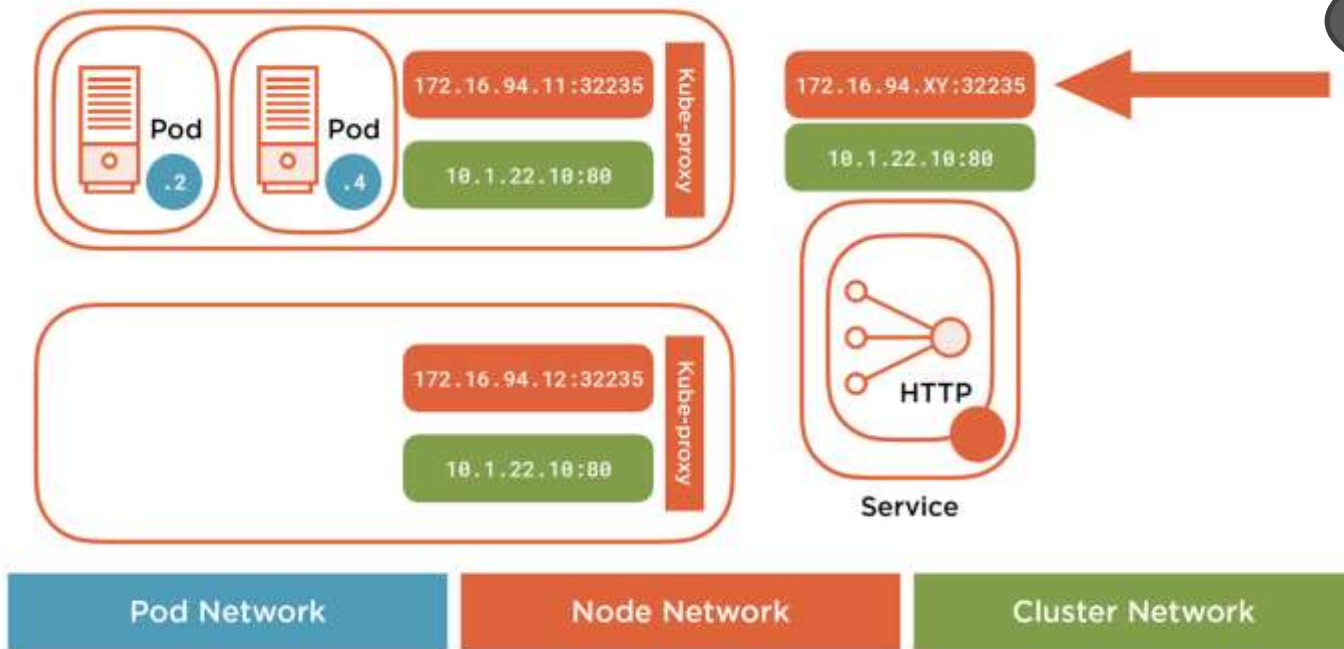




NodePort

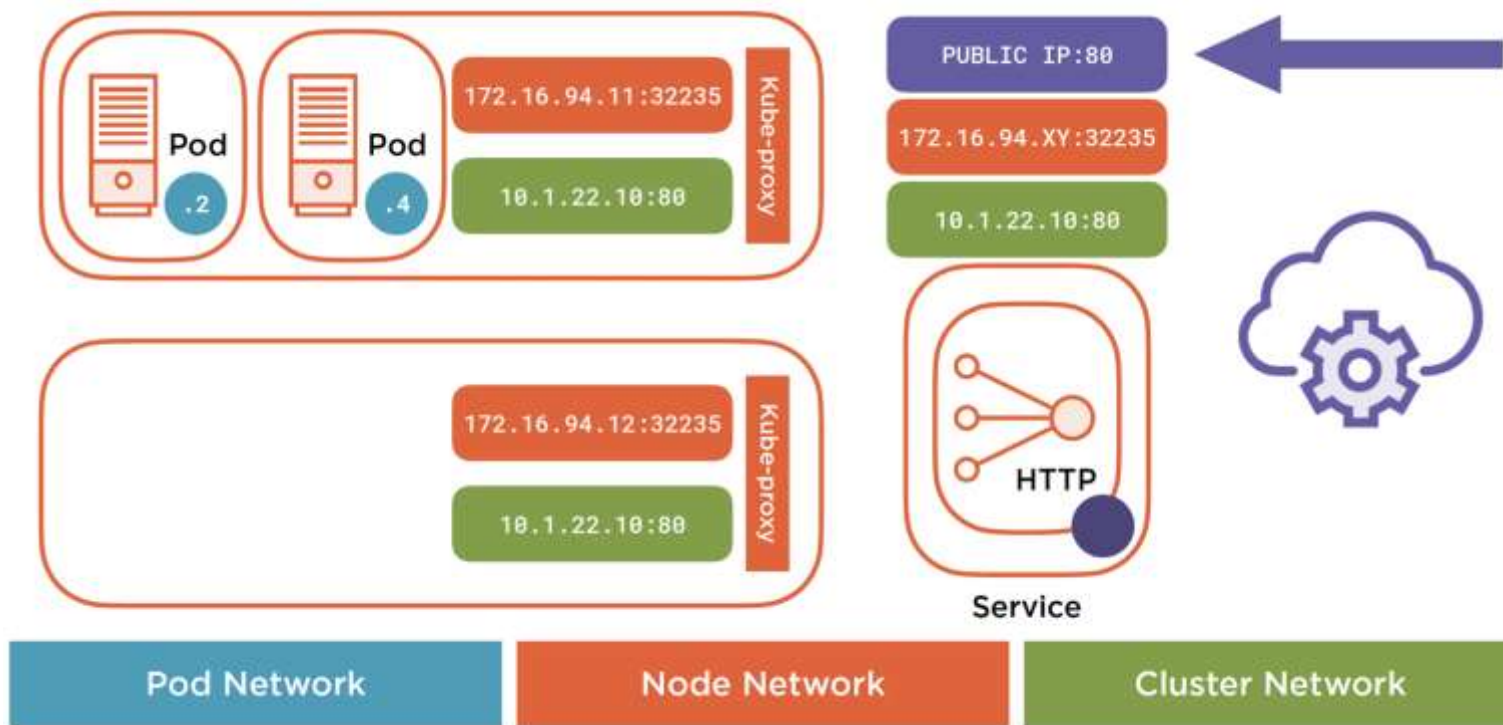
Internet

FW






LoadBalancer





LoadBalancer





Defining deployments and services

```
kind: Deployment
```

```
...
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        run: hello-world
```

```
    spec:
```

```
      containers:
```

```
...
```

```
kind: Service
```

```
...
```

```
spec:
```

```
  type: ClusterIP
```

```
  selector:
```

```
    run: hello-world
```

```
  ports:
```

```
    - port: 80
```

```
      protocol: TCP
```

```
      targetPort: 8080
```



```
kubectl create deployment hello-world --image=gcr.io/google-samples/hello-app:1.0
```

```
kubectl expose deployment hello-world --port=80 --target-port=8080 --type NodePort
```



Demo 1



Expose and access app with services

- ClusterIP
- NodePort
- LoadBalancer

