

Securing your cluster



THAOLUONG
09/2020

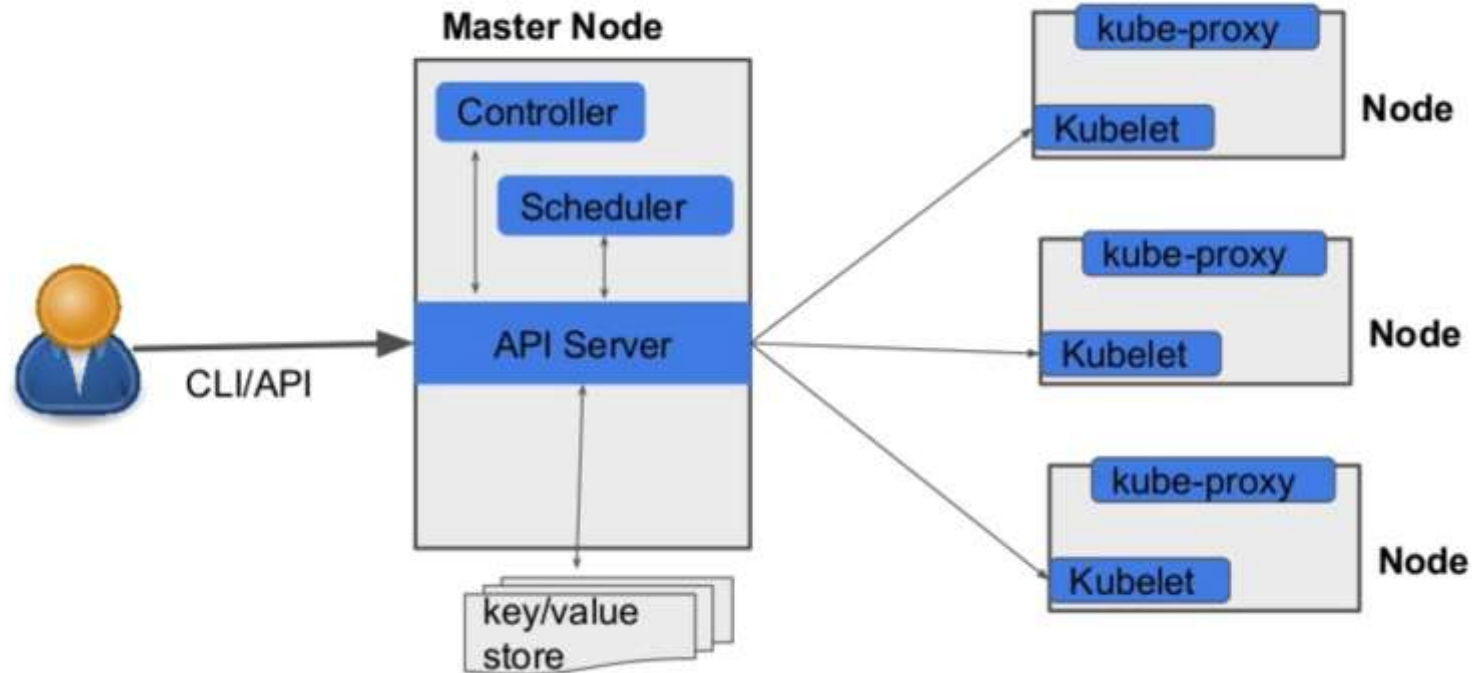


Content

- ❑ API request
- ❑ RBAC
- ❑ Kube config
- ❑ TLS certificate
- ❑ Network Policy

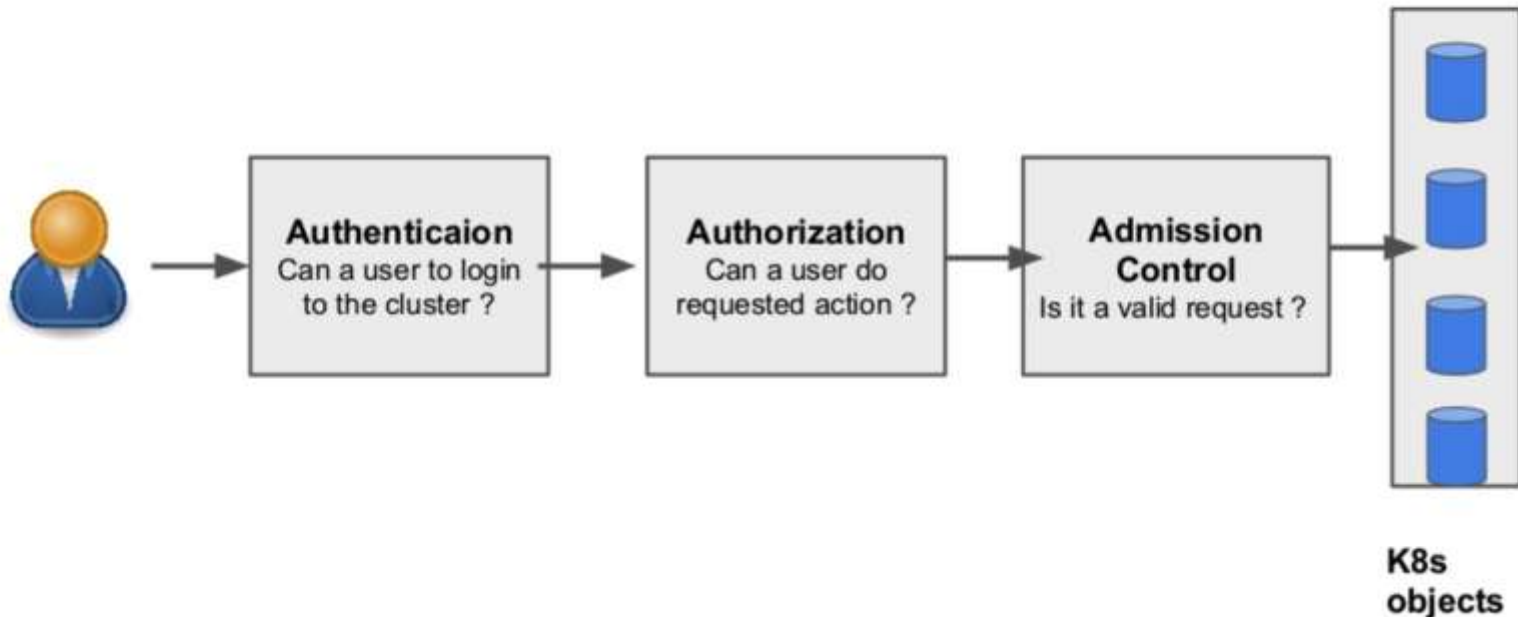


Kubernetes Architecture





Kubernetes API request





Authentication



Who can access ?

- Files – Username and Password
- Certificates
- External Authentication providers – LDAP
- Service Account



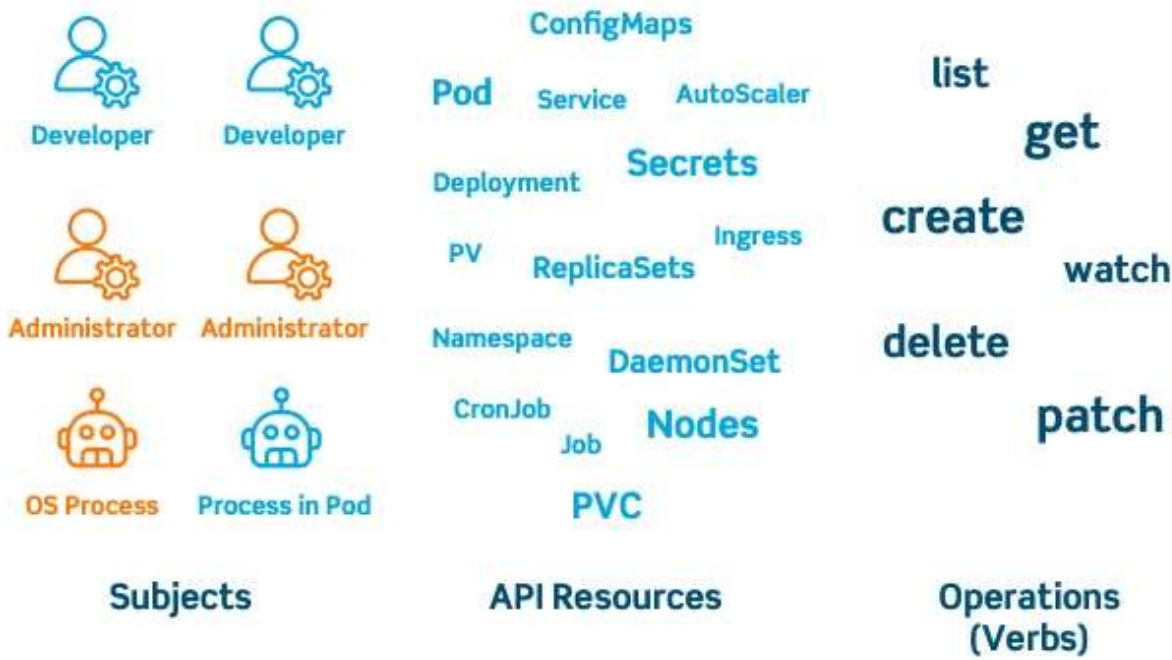
Authorization



- RBAC Authorization
- ABAC Authorization
- Node Authorization
- Webhook mode

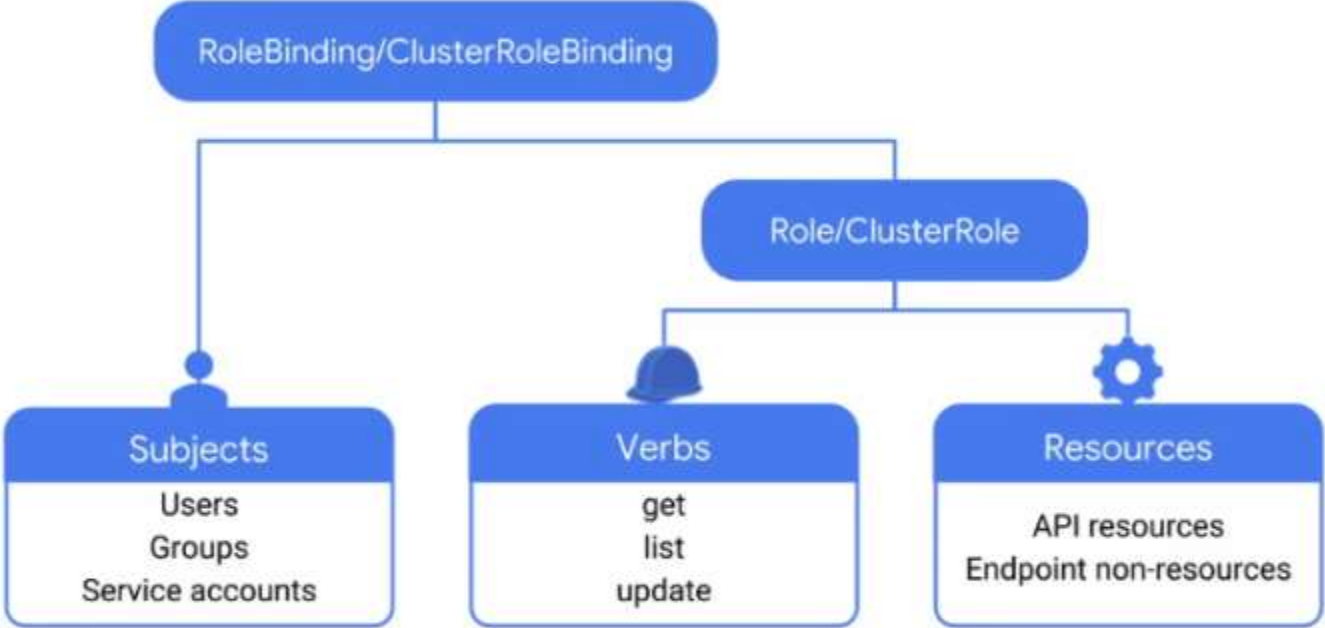


Role-Based Access Control





RBAC Structure





No User Management In Kubernetes



Expectation:

~~kubectl create user john~~

~~kubectl create group admin~~

~~kubectl add john to admin~~





How to manage user? User Plugin



- Certificate Based Authentication (X509)
- Token Based Authentication
- Basic Authentication
- OAUTH2: OIDC



Service Account

Accessing the API Server From a Pod

For example, some applications might need to know:

- The status of the cluster's nodes.
- The namespaces available.
- The Pods running in the cluster, or in a specific namespace.
- And other things like that.



RBAC - Roles

Role

“Applicable to a given namespace only.”

kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

namespace: cloudyuga

name: deployment-manager

rules:

- apiGroups: ["", "apps"]

resources: ["deployments", "replicasets", "pods"]

verbs: ["get", "list", "watch", "create", "update"]

ClusterRole

“Applicable Cluster Wide.”

kind: ClusterRole

apiVersion: rbac.authorization.k8s.io/v1

metadata:

name: deployment-manager-cluster

rules:

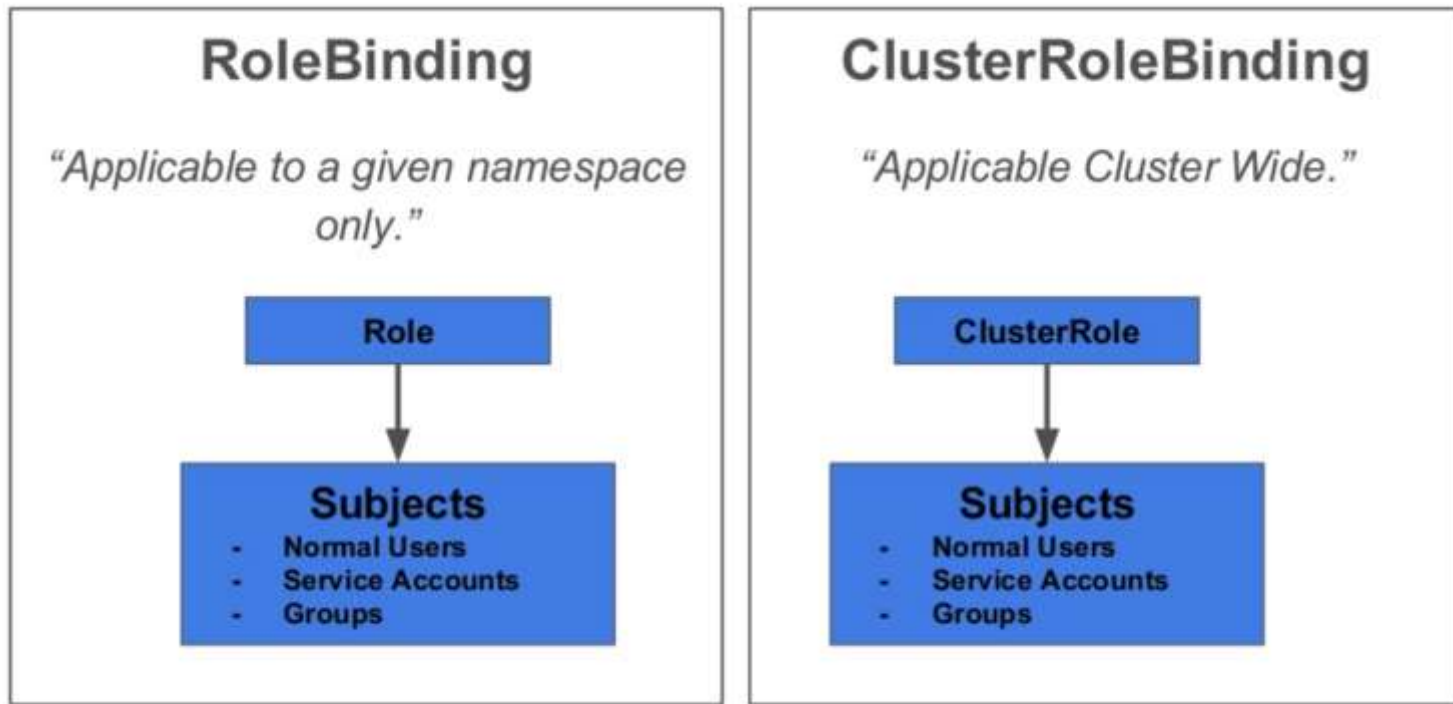
- apiGroups: ["", "apps"]

resources: ["deployments", "replicasets", "pods"]

verbs: ["get", "list", "watch", "create", "update"]



RBAC - Roles Binding





RBAC - Roles Binding

RoleBinding

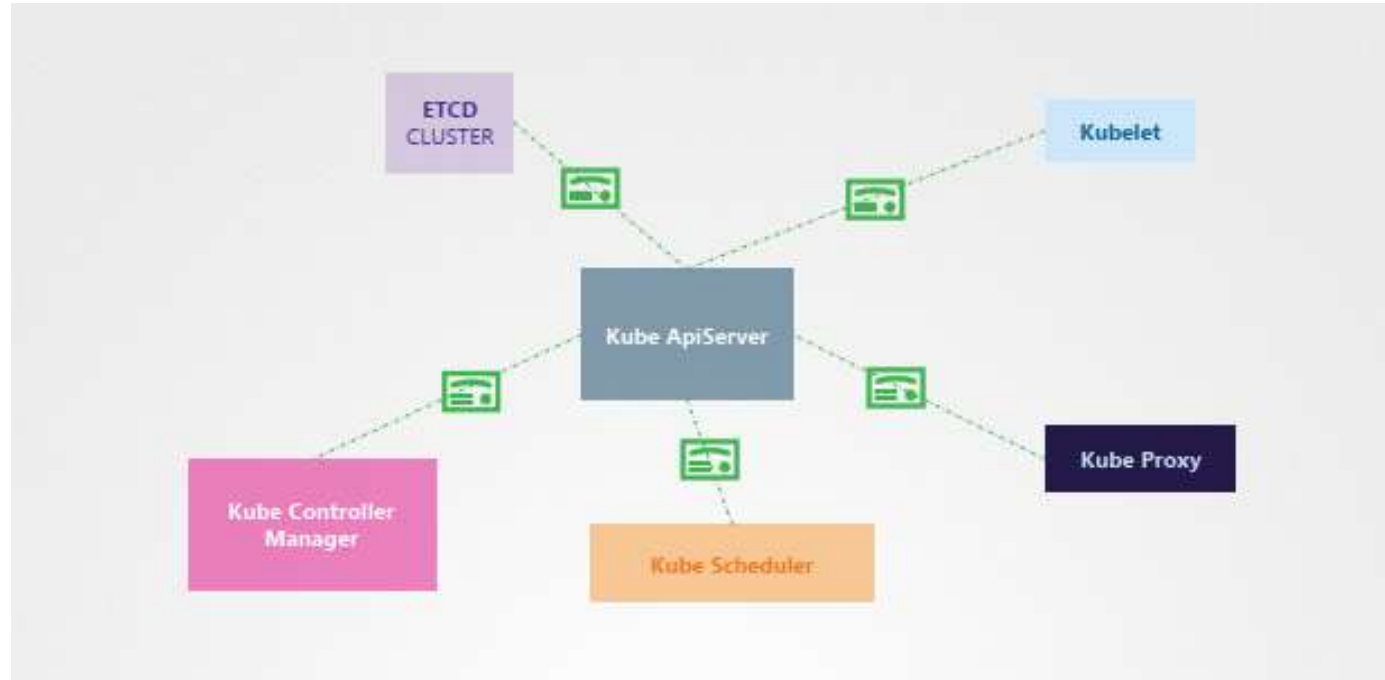
```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: deployment-manager-binding
  namespace: cloudyuga
subjects:
- kind: User
  name: nkhare
  apiGroup: "rbac.authorization.k8s.io"
roleRef:
  kind: Role
  name: deployment-manager
  apiGroup: "rbac.authorization.k8s.io"
```

ClusterRoleBinding

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: cluster-manager-binding
subjects:
- kind: User
  name: nkhare
  apiGroup: "rbac.authorization.k8s.io"
roleRef:
  kind: ClusterRole
  name: deployment-manager-cluster
  apiGroup: "rbac.authorization.k8s.io"
```

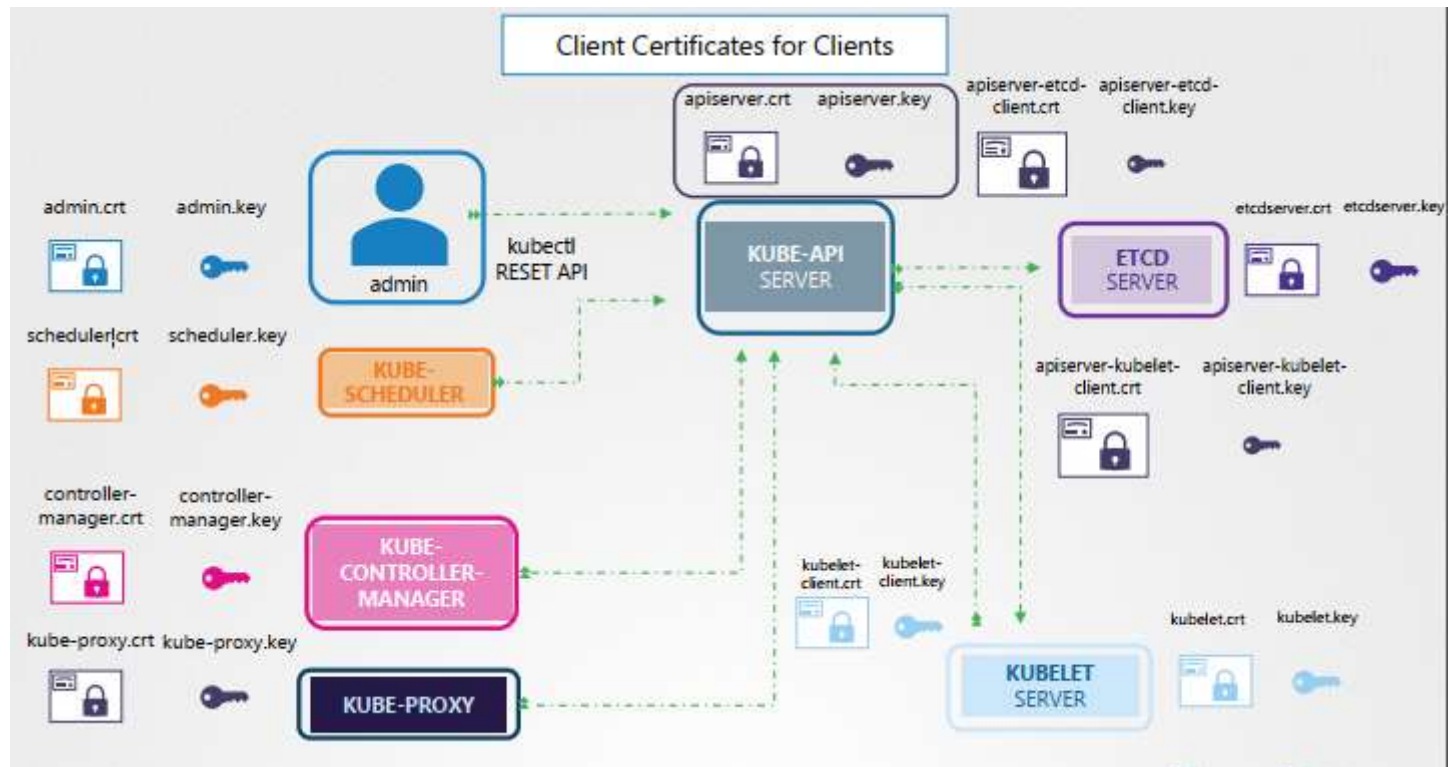


TLS certificate





TLS





Kubeconfig

```
apiVersion: v1
kind: Config

clusters:
- name: my-kube-playground
  cluster:
    certificate-authority: ca.crt
    server: https://my-kube-playground:6443

contexts:
- name: my-kube-admin@my-kube-playground
  context:
    cluster:
    user:

users:
- name: my-kube-admin
  user:
    client-certificate: admin.crt
    client-key: admin.key
```

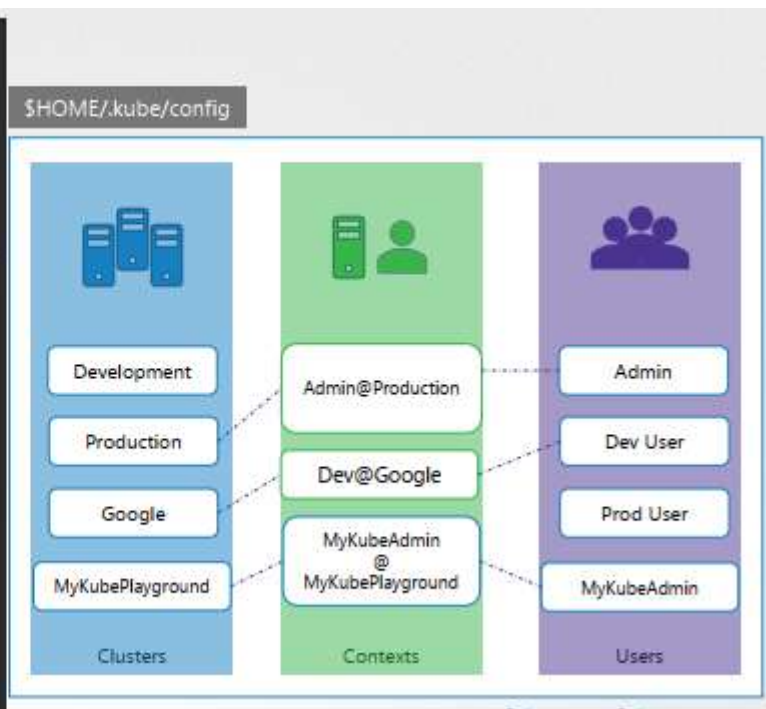




Image Security

nginx-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - name: nginx
    image: nginx
```

image: docker.io/nginx/nginx



Registry

User/
Account

Image/
Repository

gcr.io/kubernetes-e2e-test-images/dnsutils



Private repository

```
▶ docker login private-registry.io
```

```
▶ docker run private-registry.io/apps/internal-app
```

```
▶ kubectl create secret docker-registry regcred \
  --docker-server= private-registry.io \
  --docker-username= registry-user \
  --docker-password=registry-password \
  --docker-email= registry-user@org.com
```

nginx-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - name: nginx
    image: private-registry.io/apps/internal-app
  imagePullSecrets:
  - name: regcred
```



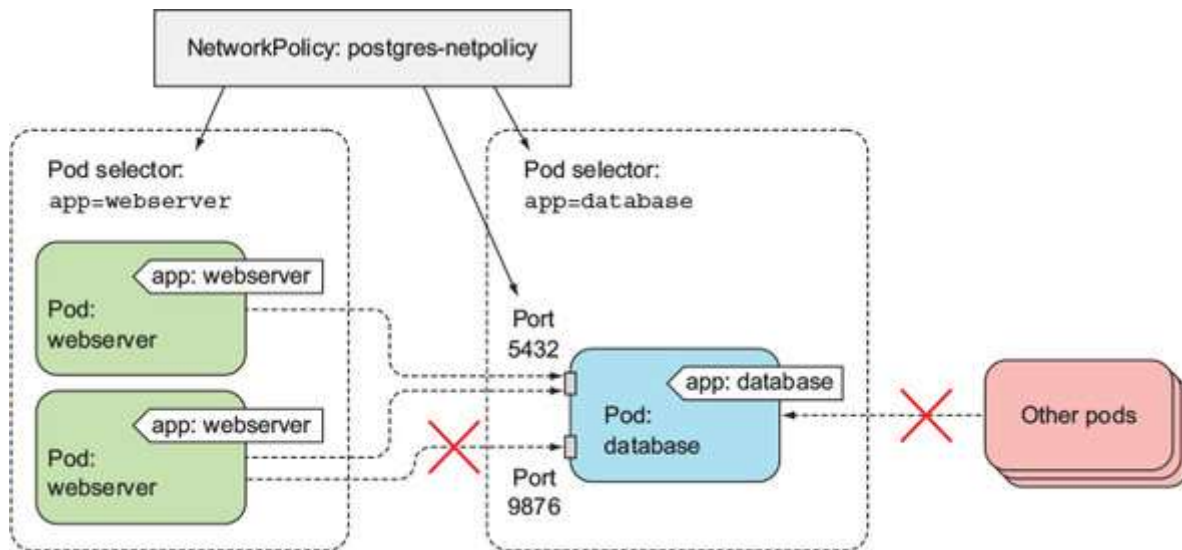
Network Policies



- By default, Kubernetes does not restrict traffic between pods running inside the cluster.
- Securing Kubernetes Cluster Networking.
- A specification of how groups of pods are allowed to communicate with each other and other network endpoints.
- Declaratively configure which pods are allowed to connect to each other.
- NetworkPolicy resources use labels to select pods and define rules which specify what traffic is allowed to the selected pods.



Network Policy Use Case



```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: postgres-netpolicy
spec:
  podSelector:
    matchLabels:
      app: database
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: webserver
      ports:
        - port: 5432
```




Network Policy Manifest

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: api-allow
spec:
  podSelector:
    matchLabels:
      app: bookstore
      role: api
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: bookstore
        - from:
            - podSelector:
                matchLabels:
                  app: inventory
```




Network Policy Use Case

Deny All



Limit Traffic



Deny From Other Namespace



Selective traffic

