

# SECURING KUBERNETES

## I. Service Account

```
## Create User
```

```
## Create Service Account
```

- Each namespace has a default ServiceAccount, named default

```
---
```

```
kubectl get sa --all-namespaces | grep default
```

```
---
```

default	default	1	5m49s
kube-node-lease	default	1	5m49s
kube-public	default	1	5m49s
kube-system	default	1	5m49s

- Let's inspect the ServiceAccount named default of the default namespace

```
---
```

```
kubectl get sa default -o yaml
```

```
---
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  creationTimestamp: "2020-07-27T00:24:14Z"
  name: default
  namespace: default
  resourceVersion: "254"
  selfLink: /api/v1/namespaces/default/serviceaccounts/default
  uid: 0468392f-1efa-442e-8583-3d96a8df7666
secrets:
- name: default-token-grng7
```

- We can see here that a Secret is provided to this ServiceAccount.

```
---
```

```
kubectl get secret default-token-grng7 -o yaml
```

```
...

apiVersion: v1
data:
  ca.crt: <>
  namespace: ZGVmYXVsdA==
  token: <>
kind: Secret
metadata:
  annotations:
    kubernetes.io/service-account.name: default
    kubernetes.io/service-account.uid: 0468392f-1efa-442e-8583-3d96a8df7666
  creationTimestamp: "2020-07-27T00:24:14Z"
  name: default-token-grng7
  namespace: default
  resourceVersion: "251"
  selfLink: /api/v1/namespaces/default/secrets/default-token-grng7
  uid: eb46dd28-0e97-48fd-9ac3-48d91f6be4db
type: kubernetes.io/service-account-token

- Decode base64

...

echo $(kubectl get secret default-token-grng7 --template={ {.data.token } } | base64 --
decode )
...

- Check how service account apply to a pod

...

https://github.com/hungtran84/k8s-cka/blob/master/d1\_managing\_cluster/03\_RBAC/pod-noserviceaccount.yaml
kubectl apply -f pod-noserviceaccount.yaml

kubectl get po/pod-default -o yaml
...

apiVersion: v1
kind: Pod
metadata:
  annotations:
    cni.projectcalico.org/podIP: 10.48.0.3/32
    kubectl.kubernetes.io/last-applied-configuration: |
```

```
{ "apiVersion": "v1", "kind": "Pod", "metadata": { "annotations": { }, "name": "pod-  
default", "namespace": "default" }, "spec": { "containers": [ { "command": [ "sleep", "10000"  
], "image": "alpine:3.9", "name": "alpine" } ] }  
  kubernetes.io/limit-ranger: 'LimitRanger plugin set: cpu request for container  
  alpine'  
  creationTimestamp: "2020-07-27T00:34:33Z"  
  name: pod-default  
  namespace: default  
  resourceVersion: "3729"  
  selfLink: /api/v1/namespaces/default/pods/pod-default  
  uid: 380d797b-ce6d-465e-b4eb-9cb804ba6eaf  
spec:  
  containers:  
  - command:  
    - sleep  
    - "10000"  
    image: alpine:3.9  
    imagePullPolicy: IfNotPresent  
    name: alpine  
    resources:  
      requests:  
        cpu: 100m  
    terminationMessagePath: /dev/termination-log  
    terminationMessagePolicy: File  
    volumeMounts:  
    - mountPath: /var/run/secrets/kubernetes.io/serviceaccount  
      name: default-token-grng7  
      readOnly: true  
    dnsPolicy: ClusterFirst  
    enableServiceLinks: true  
    nodeName: gke-cluster-1-default-pool-f4ac809e-sx18  
    priority: 0  
    restartPolicy: Always  
    schedulerName: default-scheduler  
    securityContext: { }  
    serviceAccount: default  
    serviceAccountName: default  
    terminationGracePeriodSeconds: 30  
    tolerations:  
    - effect: NoExecute  
      key: node.kubernetes.io/not-ready  
      operator: Exists  
      tolerationSeconds: 300  
    - effect: NoExecute
```

```

key: node.kubernetes.io/unreachable
operator: Exists
tolerationSeconds: 300
volumes:
- name: default-token-grng7
  secret:
    defaultMode: 420
    secretName: default-token-grng7
status:
  conditions:
  - lastProbeTime: null
    lastTransitionTime: "2020-07-27T00:34:34Z"
    status: "True"
    type: Initialized
  - lastProbeTime: null
    lastTransitionTime: "2020-07-27T00:34:36Z"
    status: "True"
    type: Ready
  - lastProbeTime: null
    lastTransitionTime: "2020-07-27T00:34:36Z"
    status: "True"
    type: ContainersReady
  - lastProbeTime: null
    lastTransitionTime: "2020-07-27T00:34:33Z"
    status: "True"
    type: PodScheduled
  containerStatuses:
  - containerID:
    docker://19f7df29614f0d9326441ea1028a51fd9632928322688bd77cac53908fb7358e
    image: alpine:3.9
    imageID: docker-
    pullable://alpine@sha256:65b3a80ebe7471beecbc090c5b2cdd0aafaeafa0715f8f12e40
    dc918a3a70e32
    lastState: {}
    name: alpine
    ready: true
    restartCount: 0
    started: true
    state:
      running:
        startedAt: "2020-07-27T00:34:35Z"
  hostIP: 10.148.0.14
  phase: Running
  podIP: 10.48.0.3
  
```

podIPs:  
- ip: 10.48.0.3  
qosClass: Burstable  
startTime: "2020-07-27T00:34:34Z"

- Important things to note here:

- \* The serviceAccountName key is set with the name of the default ServiceAccount.
- \* The information of the ServiceAccount is mounted inside the container of the Pod, through the usage of volume, in /var/run/secrets/kubernetes.io/serviceaccount

- Anonymous call of the API server

---

```
kubectl exec -ti pod-default -- sh
apk add --update curl
curl https://kubernetes/api --insecure
---
```

```
/ # curl https://kubernetes/api --insecure
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {

  },
  "status": "Failure",
  "message": "forbidden: User \"system:anonymous\" cannot get path \"/api\"",
  "reason": "Forbidden",
  "details": {

  },
  "code": 403
}/
```

- Call using the ServiceAccount token

---

```
TOKEN=$(cat /run/secrets/kubernetes.io/serviceaccount/token)
curl -H "Authorization: Bearer $TOKEN" https://kubernetes/api/v1/ --insecure
---
```

```
{
```

```
"kind": "APIResourceList",
"groupVersion": "v1",
"resources": [
  {
    "name": "bindings",
    "singularName": "",
    "namespaced": true,
    "kind": "Binding",
    "verbs": [
      "create"
    ]
  },
  {
    "name": "componentstatuses",
    "singularName": "",
    "namespaced": false,
    "kind": "ComponentStatus",
    "verbs": [
      "get",
      "list"
    ],
    "shortNames": [
      "cs"
    ]
  },
  {
    "name": "configmaps",
    "singularName": "",
    "namespaced": true,
    "kind": "ConfigMap",
    "verbs": [
      "create",
      "delete",
      "deletecollection",
      "get",
      "list",
      "patch",
      "update",
      "watch"
    ],
    "shortNames": [
      "cm"
    ],
    "storageVersionHash": "qFsyl6wFWjQ="
```

```
},  
{  
  "name": "endpoints",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "Endpoints",  
  "verbs": [  
    "create",  
    "delete",  
    "deletecollection",  
    "get",  
    "list",  
    "patch",  
    "update",  
    "watch"  
  ],  
  "shortNames": [  
    "ep"  
  ],  
  "storageVersionHash": "fWeeMqaN/OA=",  
},  
{  
  "name": "events",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "Event",  
  "verbs": [  
    "create",  
    "delete",  
    "deletecollection",  
    "get",  
    "list",  
    "patch",  
    "update",  
    "watch"  
  ],  
  "shortNames": [  
    "ev"  
  ],  
  "storageVersionHash": "r2yiGXH7wu8=",  
},  
{  
  "name": "limitranges",  
  "singularName": "",
```

```
"namespaced": true,
"kind": "LimitRange",
"verbs": [
  "create",
  "delete",
  "deletecollection",
  "get",
  "list",
  "patch",
  "update",
  "watch"
],
"shortNames": [
  "limits"
],
"storageVersionHash": "EBKMFVe6cwo="
},
{
  "name": "namespaces",
  "singularName": "",
  "namespaced": false,
  "kind": "Namespace",
  "verbs": [
    "create",
    "delete",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
    "ns"
  ],
  "storageVersionHash": "Q3oi5N2YM8M="
},
{
  "name": "namespaces/finalize",
  "singularName": "",
  "namespaced": false,
  "kind": "Namespace",
  "verbs": [
    "update"
  ]
}
```



```
},
{
  "name": "namespaces/status",
  "singularName": "",
  "namespaced": false,
  "kind": "Namespace",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "nodes",
  "singularName": "",
  "namespaced": false,
  "kind": "Node",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
    "no"
  ],
  "storageVersionHash": "XwShjMxG9Fs="
},
{
  "name": "nodes/proxy",
  "singularName": "",
  "namespaced": false,
  "kind": "NodeProxyOptions",
  "verbs": [
    "create",
    "delete",
    "get",
    "patch",
    "update"
  ]
}
```

```
},
{
  "name": "nodes/status",
  "singularName": "",
  "namespaced": false,
  "kind": "Node",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "persistentvolumeclaims",
  "singularName": "",
  "namespaced": true,
  "kind": "PersistentVolumeClaim",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
    "pvc"
  ],
  "storageVersionHash": "QWTyNDq0dC4="
},
{
  "name": "persistentvolumeclaims/status",
  "singularName": "",
  "namespaced": true,
  "kind": "PersistentVolumeClaim",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
```

```
"name": "persistentvolumes",
"singularName": "",
"namespaced": false,
"kind": "PersistentVolume",
"verbs": [
  "create",
  "delete",
  "deletecollection",
  "get",
  "list",
  "patch",
  "update",
  "watch"
],
"shortNames": [
  "pv"
],
"storageVersionHash": "HN/zwEC+JgM="
},
{
  "name": "persistentvolumes/status",
  "singularName": "",
  "namespaced": false,
  "kind": "PersistentVolume",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "pods",
  "singularName": "",
  "namespaced": true,
  "kind": "Pod",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ]
}
```

```
,
"shortNames": [
  "po"
],
"categories": [
  "all"
],
"storageVersionHash": "xPOwRZ+Yhw8="
},
{
  "name": "pods/attach",
  "singularName": "",
  "namespaced": true,
  "kind": "PodAttachOptions",
  "verbs": [
    "create",
    "get"
  ]
},
{
  "name": "pods/binding",
  "singularName": "",
  "namespaced": true,
  "kind": "Binding",
  "verbs": [
    "create"
  ]
},
{
  "name": "pods/eviction",
  "singularName": "",
  "namespaced": true,
  "group": "policy",
  "version": "v1beta1",
  "kind": "Eviction",
  "verbs": [
    "create"
  ]
},
{
  "name": "pods/exec",
  "singularName": "",
  "namespaced": true,
  "kind": "PodExecOptions",
```

```
"verbs": [  
  "create",  
  "get"  
]  
},  
{  
  "name": "pods/log",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "Pod",  
  "verbs": [  
    "get"  
  ]  
},  
{  
  "name": "pods/portforward",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "PodPortForwardOptions",  
  "verbs": [  
    "create",  
    "get"  
  ]  
},  
{  
  "name": "pods/proxy",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "PodProxyOptions",  
  "verbs": [  
    "create",  
    "delete",  
    "get",  
    "patch",  
    "update"  
  ]  
},  
{  
  "name": "pods/status",  
  "singularName": "",  
  "namespaced": true,  
  "kind": "Pod",  
  "verbs": [  
    "get",
```

```
"patch",
"update"
],
},
{
"name": "podtemplates",
"singularName": "",
"namespaced": true,
"kind": "PodTemplate",
"verbs": [
"create",
"delete",
"deletecollection",
"get",
"list",
"patch",
"update",
"watch"
],
"storageVersionHash": "LIXB2x4IFpk="
},
{
"name": "replicationcontrollers",
"singularName": "",
"namespaced": true,
"kind": "ReplicationController",
"verbs": [
"create",
"delete",
"deletecollection",
"get",
"list",
"patch",
"update",
"watch"
],
"shortNames": [
"rc"
],
"categories": [
"all"
],
"storageVersionHash": "Jond2If31h0="
},
}
```

```
{
  "name": "replicationcontrollers/scale",
  "singularName": "",
  "namespaced": true,
  "group": "autoscaling",
  "version": "v1",
  "kind": "Scale",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "replicationcontrollers/status",
  "singularName": "",
  "namespaced": true,
  "kind": "ReplicationController",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "resourcequotas",
  "singularName": "",
  "namespaced": true,
  "kind": "ResourceQuota",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
    "quota"
  ],
  "storageVersionHash": "8uhSgffRX6w="
},
```

```
{
  "name": "resourcequotas/status",
  "singularName": "",
  "namespaced": true,
  "kind": "ResourceQuota",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
},
{
  "name": "secrets",
  "singularName": "",
  "namespaced": true,
  "kind": "Secret",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "storageVersionHash": "S6u1pOWzb84="
},
{
  "name": "serviceaccounts",
  "singularName": "",
  "namespaced": true,
  "kind": "ServiceAccount",
  "verbs": [
    "create",
    "delete",
    "deletecollection",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
```



```
"sa"
],
"storageVersionHash": "pbx9ZvyFpBE="
},
{
  "name": "serviceaccounts/token",
  "singularName": "",
  "namespaced": true,
  "group": "authentication.k8s.io",
  "version": "v1",
  "kind": "TokenRequest",
  "verbs": [
    "create"
  ]
},
{
  "name": "services",
  "singularName": "",
  "namespaced": true,
  "kind": "Service",
  "verbs": [
    "create",
    "delete",
    "get",
    "list",
    "patch",
    "update",
    "watch"
  ],
  "shortNames": [
    "svc"
  ],
  "categories": [
    "all"
  ],
  "storageVersionHash": "0/CO1lhkEBI="
},
{
  "name": "services/proxy",
  "singularName": "",
  "namespaced": true,
  "kind": "ServiceProxyOptions",
  "verbs": [
    "create",
```

```
"delete",
"get",
"patch",
"update"
]
},
{
  "name": "services/status",
  "singularName": "",
  "namespaced": true,
  "kind": "Service",
  "verbs": [
    "get",
    "patch",
    "update"
  ]
}
]
```

- Let's now try something more ambitious, and use this token to list all the Pods within the default namespace

```
---
```

```
curl -H "Authorization: Bearer $TOKEN"
https://kubernetes/api/v1/namespaces/default/pods/ --insecure
```

```
---
```

```
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {

  },
  "status": "Failure",
  "message": "pods is forbidden: User \"system:serviceaccount:default:default\"
cannot list resource \"pods\" in API group \"\" in the namespace \"default\"",
  "reason": "Forbidden",
  "details": {
    "kind": "pods"
  },
  "code": 403
}/
```

## 2. Pod with service account:

```
# Create service account:
Kubect create serviceaccount demo-sa
#Create pod with sa:
https://github.com/hungtran84/k8s-cka/blob/master/d1\_managing\_cluster/03\_RBAC/pod-with-sa.yaml
kubectl apply -f pod-with-sa.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-demo-sa
spec:
  serviceAccountName: demo-sa
  containers:
  - name: alpine
    image: alpine:3.9
    command:
    - "sleep"
    - "10000"

#get the pod with sa:
Kubectl describe pod pod-demo-sa
```

## 3. Create User:

```
# Create user demouser with role permission developer only on the pod
#Generate the ssl certificate for demoUser
openssl genrsa -out demoUser.key 2048

openssl req -new -key demoUser.key -out demoUser.csr

#Get the csr file content
CSR=$(cat demoUser.csr | base64 | tr -d "\n")
```

Create CSR request

```
cat <<EOF | kubectl apply -f -
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: demoUser
spec:
  groups:
  - system:authenticated
  request: $CSR
  usages:
  - client auth
EOF
```

# list of csr:

```
kubectl get csr
```

#Approve the demoUser CSR

```
kubectl certificate approve demoUser
```

#get the certification:

```
kubectl get csr demoUser -o jsonpath='{.status.certificate}' | base64 -d >
demoUser.crt
```

#Create the role for developer

```
kubectl create role developer --verb=create --verb=get --verb=list --verb=update --
verb=delete --resource=pods
```

#Create rolebinding for developer

```
kubectl create rolebinding developer-binding-demoUser --role=developer --
user=demoUser
```

#Test with the demo user

```
k get pod --as demoUser
```

```
k get deployment --as demoUser
```

#Add new user to the kubeconfig (options):

```
kubectctl config set-credentials demoUser --client-key=demoUser.key --client-
certificate=demoUser.crt --embed-certs=true
```

```
kubectctl config set-context demoUser --cluster=kubernetes --user=demoUser
kubectctl config use-context demoUser
```

#### 4. Pulling a Container from a Private Container Registry

```
#Create secrets
```

```
kubectctl create secret docker-registry private-reg-cred \
  --docker-server=https://index.docker.io/v1/ \
  --docker-username=$DOCKERACC \
  --docker-password=$PASSWORD \
  --docker-email=$EMAIL
```

```
secret/private-reg-cred created
```

```
kubectctl apply -f deployment-private-registry.yaml
deployment.apps/hello-world-private-registry created
```

```
kubectctl get pods hello-world
```

Conditions:

Type	Status
Initialized	True
Ready	True
ContainersReady	True
PodScheduled	True

Volumes:

default-token-r5klh:

Type: Secret (a volume populated by a Secret)

SecretName: default-token-r5klh

Optional: false

QoS Class: Guaranteed

Node-Selectors: <none>

Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s  
 node.kubernetes.io/unreachable:NoExecute for 300s

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	8s	default-scheduler	Successfully assigned default/hello-world-private-registry-76f8ddc944-p2jsv to gke-cluster-1-default-pool-990b49f7-bzft
Normal	Pulling	7s	kubelet, gke-cluster-1-default-pool-990b49f7-bzft	Pulling image "votiethuy/hello:latest"
Normal	Pulled	4s	kubelet, gke-cluster-1-default-pool-990b49f7-bzft	Successfully pulled image "votiethuy/hello:latest"
Normal	Created	4s	kubelet, gke-cluster-1-default-pool-990b49f7-bzft	Created container hello-world
Normal	Started	4s	kubelet, gke-cluster-1-default-pool-990b49f7-bzft	Started container hello-world

TOKEN